

# Apple Twitter Sentiment Analysis

## Business Understanding

### Business Context

Apple is one of the most talked-about brands on social media, with millions of people sharing their opinions of its products, services, and company decisions. Understanding public sentiment from these discussions can help the company track brand perception, customer satisfaction, and market trends. Due to this, the project uses Natural Language Processing (NLP) and machine learning to classify Apple-related tweets as positive, negative, or neutral. By identifying the best-performing model, we can help the company and analysts gain valuable insights into public opinion, guiding better decision-making.

### Problem Statement

Understanding public sentiment toward Apple on Twitter is challenging due to short, informal text and varying contexts. Misclassifying sentiment can lead to inaccurate insights, affecting company's decisions. This project aims to determine the most effective sentiment analysis model by comparing traditional machine learning and deep learning approaches to achieve the highest accuracy.

## Objectives

### Main Objective

To develop an accurate sentiment analysis model for Apple-related tweets by comparing traditional machine learning and deep learning approaches.

### Specific Objectives

1. To preprocess Apple-related tweets by cleaning, tokenizing, and normalizing text data to ensure high-quality input for analysis.
2. To handle data imbalance and enhance dataset quality using techniques such as SMOTE and other resampling methods to create a well-balanced training set.
3. To develop and compare multiple sentiment classification models, including traditional machine learning such as Logistic Regression, and XGBoost and deep learning approaches such as LSTM and CNN, to identify the most effective model.
4. To evaluate model performance using appropriate metrics such as accuracy ensuring the best-performing model provides reliable sentiment insights.

## Why Machine Learning and Deep Learning?

Machine Learning (ML) and Deep Learning (DL) are well-suited for sentiment analysis due to their ability to handle large-scale text data, capture patterns in language, and generalize well across unseen data.

- **Machine Learning (ML)** models such as Logistic Regression and XGBoost are interpretable, computationally efficient, and perform well on structured text features like TF-IDF and word embeddings. These models offer quick training times and are useful for baseline comparisons.
- **Deep Learning (DL)** models like LSTM and CNN excel in understanding contextual meaning, capturing sequential dependencies, and leveraging pre-trained knowledge from large-scale corpora. These models significantly improve accuracy in sentiment classification by recognizing complex language patterns.

By combining both approaches, we can compare performance, efficiency, and scalability, ensuring the most effective model is selected for sentiment analysis.

## Success Metrics

The model's performance was evaluated using the following key metrics:

1. **Accuracy** – The percentage of correctly classified sentiments, with a target of above 70%.
2. **Overfitting Control** – The model was assessed for generalization, ensuring minimal performance gaps between training and test sets.
3. **Model Stability** – The model's consistency was tested across different subsets of data to confirm its reliability.

Success was defined as achieving these metrics while preventing overfitting and ensuring robust sentiment classification.

## Key Stakeholders

1. **Apple Inc.** – Understands public sentiment to enhance product development, marketing strategies, and customer engagement.
2. **Investors & Market Analysts** – Leverage sentiment insights to predict consumer confidence and potential stock movements.
3. **Marketing & PR Teams** – Optimize branding, crisis management, and targeted advertising based on sentiment trends.
4. **Technology Consumers & Apple Users** – Benefit from improved products, services, and

## Data Understanding

The dataset consists of **3886 tweets**, each labeled with sentiment and sentiment confidence scores.

## Sentiment Distribution

- **Neutral (3)**: 2162 tweets (Largest class)
- **Negative (1)**: 1219 tweets
- **Positive (5)**: 423 tweets

- **Not Relevant:** 82 tweets
- **Observation:** The dataset is **imbalanced**, with more neutral and negative tweets.

## Sentiment Confidence Scores

- The scores range from **0.3 to 1.0**.
- **Peaks at 0.7 and 1.0**, indicating varying label reliability.
- **High-confidence labels** can be prioritized for training to improve model accuracy.

## Tweet Length Distribution

- Most tweets are **between 100 and 140 characters**.
- A **longer tweet length** trend is observed, likely due to detailed opinions or news articles.

## Handling Missing Values

- `sentiment_gold` : **Missing in 3783 rows**, making it **unusable**.
- `_last_judgment_at` : **103 missing values**, but **not critical** for modeling.

## Duplicates

- **No duplicate tweets** found.

## Top Hashtags and Words

- **Top Hashtags:** #AAPL , #Apple , #trading , #Stocks , #iPhone6 .
- **Top Words:** "apple", "aapl", "http", "rt", indicating **frequent mentions of Apple products, financial discussions, and retweets**.

```
In [ ]: # Import the necessary libraries
# General Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import re
import warnings
warnings.filterwarnings('ignore')

# NLP Libraries
import nltk
from nltk.tokenize import word_tokenize
from nltk import pos_tag
from nltk.corpus import stopwords
import contractions
from textblob import TextBlob
from wordcloud import WordCloud
from gensim.models import Word2Vec

# TensorFlow & Keras
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import (
    Embedding, LSTM, Bidirectional, Conv1D, MaxPooling1D,
    Flatten, Dense, Dropout, BatchNormalization,
    GlobalAveragePooling1D, Input
)
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.optimizers.schedules import ExponentialDecay
from tensorflow.keras.regularizers import l2

# Scikit-Learn & ML models
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from sklearn.metrics import accuracy_score, classification_report

# XGBoost
from xgboost import XGBClassifier

# Imbalanced data handling
from imblearn.over_sampling import SMOTE
from imblearn.combine import SMOTETomek

from collections import Counter
```

```
In [ ]: # Load the dataset
df = pd.read_csv("Apple-Twitter-Sentiment-DFE.csv", encoding="ISO-8859-1")
```

```
In [ ]: # Display the first few rows
df.head()
```

Out[4]:

	_unit_id	_golden	_unit_state	_trusted_judgments	_last_judgment_at	sentiment	sentiment:coi
0	623495513	True	golden	10	NaN	3	
1	623495514	True	golden	12	NaN	3	
2	623495515	True	golden	10	NaN	3	
3	623495516	True	golden	17	NaN	3	
4	623495517	False	finalized	3	12/12/14 12:14	3	



```
In [ ]: # Displaying the last 5 rows
df.tail()
```

Out[5]:

	_unit_id	_golden	_unit_state	_trusted_judgments	_last_judgment_at	sentiment	sentiment
--	----------	---------	-------------	--------------------	-------------------	-----------	-----------

3881	623499442	True	golden	13	NaN	3	
------	-----------	------	--------	----	-----	---	--

3882	623499450	True	golden	16	NaN	3	
------	-----------	------	--------	----	-----	---	--

3883	623499486	True	golden	14	NaN	5	
------	-----------	------	--------	----	-----	---	--

3884	623499514	True	golden	13	NaN	1	
------	-----------	------	--------	----	-----	---	--

3885	623517290	True	golden	17	NaN	5	
------	-----------	------	--------	----	-----	---	--



```
In [ ]: # Check dataset shape
print("Shape:", df.shape)
```

Shape: (3886, 12)

```
In [ ]: # Check the unique values
df.nunique()
```

```
Out[7]:
```

	0
<b>_unit_id</b>	3886
<b>_golden</b>	2
<b>_unit_state</b>	2
<b>_trusted_judgments</b>	19
<b>_last_judgment_at</b>	388
<b>sentiment</b>	4
<b>sentiment:confidence</b>	654
<b>date</b>	3795
<b>id</b>	3
<b>query</b>	1
<b>sentiment_gold</b>	9
<b>text</b>	3219

**dtype:** int64


```
In [ ]: df['sentiment'].value_counts()
```

```
Out[8]:
```

	count
<b>sentiment</b>	
<b>3</b>	2162
<b>1</b>	1219
<b>5</b>	423
<b>not_relevant</b>	82

**dtype:** int64

```
In [ ]: # Count sentiment labels
df['sentiment'].value_counts().plot(kind='bar', title="Sentiment Distribution")
plt.xlabel("Sentiment")
plt.ylabel("Count")
plt.show()
```





```
In [ ]: # Set column width to display full tweets
pd.options.display.max_colwidth = None

# Display sample tweets for each sentiment category
for sentiment_value in df['sentiment'].unique():
    print(f"Sentiment: {sentiment_value}")
    print(df[df['sentiment'] == sentiment_value]['text'].sample(3, random_state=42))
    print("\n" + "="*80 + "\n")
```

Sentiment: 3

Photographing the White House Christmas Decorations With an iPhone 6 by @BrooksKraftFoto @apple <http://t.co/lPDqbJqnV5> (<http://t.co/lPDqbJqnV5>)

#Apple Wants To Make Your Commute Much Easier, According To This New Patent #aapl <http://t.co/fKMNHcmwJU> (<http://t.co/fKMNHcmwJU>) <http://t.co/wdqAzQowt3> (<http://t.co/wdqAzQowt3>)

RT @OneRepublic: Studio at 45,000 ft. One outlet, 4 computers. @Apple we need the batteries of the future Now!!!! <http://t.co/astp9x6KET> (<http://t.co/astp9x6KET>)

=====

Sentiment: 5

@MhDaDon @Apple def gotta have it, I don't even like watches fun..fun nights..Post birthday celebration of rfrancoben and @apple. <http://t.co/mARHLxgV0F> (<http://t.co/mARHLxgV0F>)

I'm really enjoying GarageBand. @apple #GarageBand

=====

Sentiment: 1

RT @OneRepublic: Studio at 45,000 ft. One outlet, 4 computers. @Apple we need the batteries of the future Now!!!! <http://t.co/astp9x6KET> (<http://t.co/astp9x6KET>)

How is 'never' interpreted as 'ask me again annoyingly soon' when iOS apps ask to be rated? @apple tell your devs never means NEVER

Thanks @apple for changing yet another fuck into duck...Thanks.

=====

Sentiment: not\_relevant

@sex  
tsatan @Applebees @Apple APPLEBEES FAVED OMG  
@Apple John Cantlie has been a prisoner of ISIS for 739 days, show you have not abandoned him. Sign <https://t.co/WTn4fuiJ0P> (<https://t.co/WTn4fuiJ0P>)  
#Samsung Sale Puts Spotlight On The Buyer, #Corni  
ng #GLW #AAPL #SSNLF <http://t.co/oFQx1Go5eL> (<http://t.co/oFQx1Go5eL>)

=====

- Sentiment 3 (Neutral/Mixed): News articles, patents, and general discussions without strong emotion.
- Sentiment 5 (Positive): Praising Apple products, expressing excitement.
- Sentiment 1 (Negative): Complaints, frustrations, sarcastic remarks.

```
In [ ]: # Check for missing values
df.isnull().sum()
```

```
Out[11]:
```


	0
<hr/>	
_unit_id	0
_golden	0
_unit_state	0
_trusted_judgments	0
_last_judgment_at	103
sentiment	0
sentiment:confidence	0
date	0
id	0
query	0
sentiment_gold	3783
text	0

**dtype:** int64

```
In [ ]: # Duplicates
df.duplicated().sum()
```


```
Out[12]: 0
```

```
In [ ]: # Distribution of Sentiment Confidence Scores
sns.histplot(df['sentiment:confidence'], bins=20, kde=True)
plt.xlabel("Sentiment Confidence")
plt.ylabel("Count")
plt.title("Distribution of Sentiment Confidence Scores")
plt.show()
```



```
In [ ]: # Tweet Length Distribution
df["tweet_length"] = df["text"].str.len()

sns.histplot(df["tweet_length"], bins=20, kde=True)
plt.xlabel("Tweet Length")
plt.ylabel("Count")
plt.title("Distribution of Tweet Lengths")
plt.show()
```



```
In [ ]: # Common Words & Hashtags

# Join all tweets into one string
all_text = " ".join(df["text"].dropna())

# Extract hashtags
hashtags = re.findall(r"#\w+", all_text)
hashtag_counts = Counter(hashtags).most_common(10)

# Extract words (excluding stopwords & special characters)
words = re.findall(r"\b\w+\b", all_text.lower())
word_counts = Counter(words).most_common(10)

print("Top 10 Hashtags:", hashtag_counts)
print("Top 10 Words:", word_counts)
```

Top 10 Hashtags: [('#AAPL', 569), ('#aapl', 466), ('#Apple', 251), ('#DieIn', 152), ('#iPhone', 64), ('#iPhone6', 57), ('#apple', 55), ('#December', 54), ('#trading', 48), ('#Stocks', 39)]

Top 10 Words: [('apple', 3957), ('t', 2597), ('co', 2324), ('http', 2269), ('the', 1701), ('aapl', 1385), ('to', 1053), ('in', 870), ('is', 868), ('rt', 848)]

## Data Cleaning/Text Cleaning

Data Cleaning involved the following

### 1. Lowercasing

-Converting all text to lowercase to ensure uniformity.

### 2. Removing URLs

-Eliminating links (http://..., www...) as they don't contribute to sentiment analysis.

### 3. Removing Mentions

-Deleting @username to focus on tweet content rather than tagged users.

### 4. Removing Hashtags

-Stripping hashtags (#Apple, #iPhone) as they were not needed for text analysis.

### 5. Removing Special Characters

-Keeping only alphanumeric text and spaces, removing punctuation or symbols.

### 6. Removing Extra Spaces

-Ensuring there were no unnecessary spaces between words.

## 7. Removing Stopwords

-Filtering common words like "the", "is", "and" while keeping negations (not, no, never) to preserve meaning.

## 8. Handling Duplicates

-Removing duplicate tweets to avoid bias in the dataset.

```
In [ ]: df.head(20)
```



Out[16]:

	_unit_id	_golden	_unit_state	_trusted_judgments	_last_judgment_at	sentiment	sentiment:
0	623495513	True	golden	10	NaN	3	
1	623495514	True	golden	12	NaN	3	
2	623495515	True	golden	10	NaN	3	
3	623495516	True	golden	17	NaN	3	
4	623495517	False	finalized	3	12/12/14 12:14	3	
5	623495518	True	golden	13	NaN	3	
6	623495519	True	golden	13	NaN	5	
7	623495520	True	golden	9	NaN	5	
8	623495521	True	golden	15	NaN	3	
9	623495522	False	finalized	3	12/12/14 0:52	3	
10	623495523	True	golden	12	NaN	1	

	<b>_unit_id</b>	<b>_golden</b>	<b>_unit_state</b>	<b>_trusted_judgments</b>	<b>_last_judgment_at</b>	<b>sentiment</b>	<b>sentiment:</b>
<b>11</b>	623495524	True	golden	9	NaN	3	
<b>12</b>	623495525	True	golden	11	NaN	3	
<b>13</b>	623495526	False	finalized	3	12/12/14 21:38	5	
<b>14</b>	623495527	True	golden	17	NaN	1	
<b>15</b>	623495528	False	finalized	6	12/12/14 15:50	3	
<b>16</b>	623495529	True	golden	16	NaN	1	
<b>17</b>	623495530	False	finalized	3	12/12/14 3:38	not_relevant	
<b>18</b>	623495531	False	finalized	3	12/12/14 4:59	3	
<b>19</b>	623495532	False	finalized	3	12/12/14 20:59	3	

```
In [ ]: #Extracting just the important columns needed for this analysis
#that is, sentiment and text
```

```
df = df[["date" , "sentiment:confidence", 'sentiment', 'text']]
df.head(10)
```

Out[17]:

	date	sentiment:confidence	sentiment	text
0	Mon Dec 01 19:30:03 +0000 2014	0.6264	3	#AAPL:The 10 best Steve Jobs emails ever...http://t.co/82G1kL94tx
1	Mon Dec 01 19:43:51 +0000 2014	0.8129	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flash Crash Today \$AAPL #aapl\nhttp://t.co/hGFcjYa0E9
2	Mon Dec 01 19:50:28 +0000 2014	1.0000	3	My cat only chews @apple cords. Such an #AppleSnob.
3	Mon Dec 01 20:26:34 +0000 2014	0.5848	3	I agree with @jimcramer that the #IndividualInvestor should own not trade #Apple #AAPL, it's extended so today's pullback is good to see
4	Mon Dec 01 20:29:33 +0000 2014	0.6474	3	Nobody expects the Spanish Inquisition #AAPL
5	Mon Dec 01 20:30:03 +0000 2014	0.5975	3	#AAPL:5 Rocket Stocks to Buy for December Gains: Apple and More...http://t.co/eG5XhXdLLS
6	Mon Dec 01 20:32:45 +0000 2014	0.8468	5	Top 3 all @Apple #tablets. Damn right! http://t.co/RJiGn2JUuB
7	Mon Dec 01 20:34:31 +0000 2014	0.6736	5	CNBCTV: #Apple's margins better than expected? #aapl http://t.co/7geVrtOGLK
8	Mon Dec 01 20:36:47 +0000 2014	0.7997	3	Apple Inc. Flash Crash: What You Need to Know http://t.co/YJlgtifdAj #AAPL
9	Mon Dec 01 20:45:03 +0000 2014	0.6360	3	#AAPL:This Presentation Shows What Makes The World's Biggest Tech Companies ...http://t.co/qIH9PqSoSd

```
In [ ]: #checking null values
df.isnull().sum()
```


Out[18]:

```
0
date 0
sentiment:confidence 0
sentiment 0
text 0
```

dtype: int64

```
In [ ]: #checking value count in sentiment column
print(df.sentiment.value_counts())
print(df.sentiment.value_counts().plot(kind='pie', title="Sentiment Distribution")
```

```
sentiment
3          2162
1          1219
5           423
not_relevant    82
Name: count, dtype: int64
Axes(0.22375,0.11;0.5775x0.77)
```



```
In [ ]: #removing unnecessary row not_relevant because it does not contribute to the anal
df = df[df['sentiment'] != 'not_relevant']
print(df.sentiment.unique())
```

```
['3' '5' '1']
```

```
In [ ]: #checking duplicates  
print(df.duplicated().sum())  
duplicates = df[df.duplicated(subset=["text"], keep=False)]  
print(duplicates)
```

```

4
                                date  sentiment:confidence  sentiment  \
29    Tue Dec 02 00:15:26 +0000 2014                1.0000        3
32    Tue Dec 02 00:16:27 +0000 2014                0.6604        3
34    Tue Dec 02 00:18:59 +0000 2014                0.6515        3
38    Tue Dec 02 00:24:26 +0000 2014                1.0000        3
42    Tue Dec 02 00:27:36 +0000 2014                1.0000        3
...
3852  Tue Dec 09 21:12:55 +0000 2014                0.7325        3
3854  Tue Dec 09 21:14:04 +0000 2014                1.0000        1
3855  Tue Dec 09 21:17:24 +0000 2014                0.6785        1
3878  Tue Dec 09 21:24:22 +0000 2014                0.6839        5
3885  Tue Dec 09 09:01:25 +0000 2014                0.8938        5

```

```
text
```

```

29                                RT @thehill: Justice Department cites 18th century federa
l law to get @Apple to unlock iPhones: http://t.co/Eth0QpAIom (http://t.co/Eth0QpAIom)
32                                RT @thehill: Justice Department cites 18th century federa
l law to get @Apple to unlock iPhones: http://t.co/Eth0QpAIom (http://t.co/Eth0QpAIom)
34                                RT @thehill: Justice Department cites 18th century federa
l law to get @Apple to unlock iPhones: http://t.co/Eth0QpAIom (http://t.co/Eth0QpAIom)
38                                RT @thehill: Justice Department cites 18th century federa
l law to get @Apple to unlock iPhones: http://t.co/Eth0QpAIom (http://t.co/Eth0QpAIom)
42                                RT @thehill: Justice Department cites 18th century federa
l law to get @Apple to unlock iPhones: http://t.co/Eth0QpAIom (http://t.co/Eth0QpAIom)
...
...
3852  RT @TeamCavuto: Protesters stage #DieIn protests in @Apple store in NY
C... Is it me, or is this anger misplaced? RETWEET if you agree.
3854                                RT @Ecofantasy: Thinking of upgrading to #Yosemite? Think twice h
ttp://t.co/dU0Mpaw5Ri (http://t.co/dU0Mpaw5Ri) It's not for everyone. RT #ASMSG
@Apple
3855                                RT @Ecofantasy: Thinking of upgrading to #Yosemite? Think twice h
ttp://t.co/dU0Mpaw5Ri (http://t.co/dU0Mpaw5Ri) It's not for everyone. RT #ASMSG
@Apple
3878                                RT @shannonmmiller: Love the @Apple is sup
porting #HourOfCode with workshops! :) http://t.co/WP8D0FNjNu (http://t.co/WP8D0FNjNu)
3885                                RT @SwiftKey: We're so excited to be named to @Apple's 'App Stor
e Best of 2014' list this year! http://t.co/d7qlmti4Uf (http://t.co/d7qlmti4Uf)
#Apple

```

```
[730 rows x 4 columns]
```

```
In [ ]: #checking duplicates
df[df.duplicated()]
```

```
Out[22]:
```

	date	sentiment:confidence	sentiment	text
1437	Thu Dec 04 20:39:48 +0000 2014	1.0	3	RT @OneRepublic: Studio at 45,000 ft. One outlet, 4 computers. @Apple we need the batteries of the future NoW!!!! <a href="http://t.co/astp9x6KET">http://t.co/astp9x6KET</a>
1445	Thu Dec 04 20:39:55 +0000 2014	1.0	3	RT @OneRepublic: Studio at 45,000 ft. One outlet, 4 computers. @Apple we need the batteries of the future NoW!!!! <a href="http://t.co/astp9x6KET">http://t.co/astp9x6KET</a>
1449	Thu Dec 04 20:39:58 +0000 2014	1.0	3	RT @OneRepublic: Studio at 45,000 ft. One outlet, 4 computers. @Apple we need the batteries of the future NoW!!!! <a href="http://t.co/astp9x6KET">http://t.co/astp9x6KET</a>
2511	Sat Dec 06 18:46:30 +0000 2014	1.0	1	NO @apple NO! When I make an I phone Album I WANT IT TO STAY ON PHONE, not be removed when camera roll cleared.. GET IT TOGETHER!

***There were are no duplicates just retweets***

```
In [ ]: #convert date to date_time format
#convert sentiment to integer
print(df.info())

df['date'] = pd.to_datetime(df['date'], errors='coerce')
df['sentiment'] = df['sentiment'].fillna(99).astype(int)
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3804 entries, 0 to 3885
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                   3804 non-null   object
1   sentiment:confidence   3804 non-null   float64
2   sentiment              3804 non-null   object
3   text                   3804 non-null   object
dtypes: float64(1), object(3)
memory usage: 148.6+ KB
None
```

***Defining the text cleaning function***

```
In [ ]: # Ensuring stopwords are available
nltk.download("stopwords")
nltk.download("wordnet")
stop_words = set(stopwords.words("english")) - {"not", "no", "never"} # Keep neg

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
```

```
In [ ]: # Define text cleaning function
def clean_text(text):
    # Convert to lowercase. Converting all text to lowercase to ensure uniformity
    text = text.lower().strip()
    # Expand contractions
    text = contractions.fix(text)
    # Eliminating links (http://..., www...) as they don't contribute to sentiment
    text = re.sub(r"http\S+|www\S+", "", text)
    # Deleting @username to focus on tweet content rather than tagged users.
    text = re.sub(r"@w+", "", text)
    # Stripping hashtags (#Apple, #iPhone) as they were not be needed for text analysis
    text = re.sub(r"#[A-Za-z0-9]+", "", text)
    # Keeping only alphanumeric text and spaces, removing punctuation or symbols.
    text = re.sub(r"[^A-Za-z0-9 ]+", "", text)
    # Ensuring there were no unnecessary spaces between words.
    text = re.sub(r"\s+", " ", text)
    # Filtering common words like "the", "is", "and" while keeping negations (not,
    words = text.split()
    words = [word for word in words if word not in stop_words] # Remove stopwords
    return " ".join(words)
```

```
In [ ]: # Apply cleaning to tweets
df["cleaned_text"] = df["text"].apply(clean_text)
```



```
In [ ]: df.head(50)
```

Out[28]:

	date	sentiment:confidence	sentiment	text	cleaned_text
0	2014-12-01 19:30:03+00:00	0.6264	3	#AAPL:The 10 best Steve Jobs emails ever...http://t.co/82G1kL94tx	10 best steve jobs emails ever
1	2014-12-01 19:43:51+00:00	0.8129	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flash Crash Today \$AAPL #aapl\nhttp://t.co/hGFcjYa0E9	rt aapl stock miniflash crash today aapl
2	2014-12-01 19:50:28+00:00	1.0000	3	My cat only chews @apple cords. Such an #AppleSnob.	cat chews cords
3	2014-12-01 20:26:34+00:00	0.5848	3	I agree with @jimcramer that the #IndividualInvestor should own not trade #Apple #AAPL, it's extended so today's pullback is good to see	agree not trade extended todays pullback good see
				nobody	

## Feature Engineering

```
In [ ]: # Compute word count, character count, and average word length
df["word_count"] = df["cleaned_text"].apply(lambda x: len(x.split()))
df["char_count"] = df["cleaned_text"].apply(len)
df["avg_word_length"] = df["char_count"] / df["word_count"]

df[["cleaned_text", "word_count", "char_count", "avg_word_length"]].head(10)
```

Out[29]:

	cleaned_text	word_count	char_count	avg_word_length
0	10 best steve jobs emails ever	6	30	5.000000
1	rt aapl stock miniflash crash today aapl	7	40	5.714286
2	cat chews cords	3	15	5.000000
3	agree not trade extended todays pullback good see	8	49	6.125000
4	nobody expects spanish inquisition	4	34	8.500000
5	5 rocket stocks buy december gains apple	7	40	5.714286
6	top 3 damn right	4	16	4.000000
7	cnbctv margins better expected	4	30	7.500000
8	apple inc flash crash need know	6	31	5.166667
9	presentation shows makes worlds biggest tech companies	7	54	7.714286

```
In [ ]: filtered_df = df[df["word_count"] > 16] # Filter rows where avg_word_length > 15
print(filtered_df[["cleaned_text", "word_count", "sentiment"]].head(10)) # Display
```

```
cleaned_text \
69          force people use vpn built ios8 button not work ffs like want
apples nsa data collection service
98      hate ios 8 capitalizes random words like not want give emphasis stupid w
ord tha sentence get self together
394          hey love ya lowfi hold music awful would prefer hear tips usi
ng apple gear better use hold time
1164          11593 dec1 64 one crazy minute w 67m shares ms downgrad
e market weight amp trim stock 4 3
1324  could really kick ass iphone 6 battery sucks moldy dick tuesday night wo
rst shit ever last 4 fucking hours
1388      spent 6000 eur apple iphone 6 camera no longer workstold got water iti
not unacceptable customer service
1391  rt spent 6000 eur apple iphone 6 camera no longer workstold got water iti
not unacceptable customer service
2271          mark words wild away iphone 5c bring back 4 iphone 5s ultima
te form factor welcome iphone mini
2313          cgk laptop prob today local useless tech support useless 1 hr g
enius bar useless buy pc next time
2513      hell thought let us put volume display front video absolutely dumb m
iss video every time adjust volume
```

	word_count	sentiment
69	17	1
98	18	1
394	18	1
1164	18	3
1324	19	1
1388	17	1
1391	18	1
2271	17	3
2313	18	1
2513	17	1

```
In [ ]: # Compute subjectivity using TextBlob
df["subjectivity"] = df["cleaned_text"].apply(lambda x: TextBlob(x).sentiment.subjectivity)

# Display the first few rows to check the computed subjectivity scores
df[["cleaned_text", "subjectivity", "sentiment", "sentiment:confidence"]].head(10)
```

Out[31]:

	cleaned_text	subjectivity	sentiment	sentiment:confidence
0	10 best steve jobs emails ever	0.300000	3	0.6264
1	rt aapl stock miniflash crash today aapl	0.000000	3	0.8129
2	cat chews cords	0.000000	3	1.0000
3	agree not trade extended todays pullback good see	0.600000	3	0.5848
4	nobody expects spanish inquisition	0.000000	3	0.6474
5	5 rocket stocks buy december gains apple	0.000000	3	0.5975
6	top 3 damn right	0.517857	5	0.8468
7	cnbctv margins better expected	0.450000	5	0.6736
8	apple inc flash crash need know	0.000000	3	0.7997
9	presentation shows makes worlds biggest tech companies	0.000000	3	0.6360

**Observations:**

Subjectivity Scores:

- Values range from 0 (objective) to 1 (highly subjective).
- Some tweets have 0.0, indicating factual statements.
- Others, like "agree not trade extended todays pullback good see", have higher subjectivity (0.6), meaning they express opinions rather than facts.

## Exploratory Data Analysis (EDA)

### 1. Univariate Analysis

```
In [ ]: #pip install --upgrade pillow wordcloud
```



[illegible]

[illegible]

Overall Word cloud:


- Positive Word Cloud:

- Negative Word Cloud:

- Neutral Word Cloud:


- A more neutral cloud focusing on keywords like "studio," "outlet," "computers," and "batteries." This indicates general discussions about Apple products without a strong emotional tone.

```
In [ ]: # Character Length Distribution
plt.figure(figsize=(7, 4))
sns.histplot(df["char_count"], bins=30, kde=True, color="blue")
plt.title("Distribution of Tweet Length", fontsize=14)
plt.xlabel("Tweet Length (Characters)")
plt.ylabel("Frequency")
plt.show()
```



- It ranges between 60- 70 characters


```
In [ ]: # Histogram for average word length
plt.figure(figsize=(10,5))
sns.histplot(df["avg_word_length"], bins=20, kde=True)
plt.xlabel("Average Word Length")
plt.ylabel("Frequency")
plt.title("Distribution of Average Word Length in Tweets")
plt.show()
```



- The average word length is mostly around 5-7 characters, indicating that most words in the dataset are relatively short.



```
In [ ]: # Histogram for subjectivity scores
plt.figure(figsize=(10,5))
sns.histplot(df["subjectivity"], bins=20, kde=True)
plt.xlabel("Subjectivity Score")
plt.ylabel("Frequency")
plt.title("Distribution of Subjectivity Scores in Tweets")
plt.show()
```



This plot reveals that:

- Most tweets are objective → Subjectivity scores close to 0
- Only a smaller portion are strongly opinionated → Scores near 1
- That suggests many tweets are news, updates, or factual statements rather than personal opinions—useful insight for understanding tone on social media


## 2. Bivariate Analysis

```
In [ ]: #Sentiment distribution over time
#Group by date and calculate the mean sentiment
sentiment_trend = df.groupby(df['date'].dt.date)['sentiment'].mean()

#Plot
plt.figure(figsize=(10, 5))
sentiment_trend.plot(marker="o", color="royalblue")

#Labels and title
plt.xlabel("Date")
plt.ylabel("Average Sentiment Score")
plt.title("Sentiment Trend Over Time")

#Show plot
plt.xticks(rotation=45)
plt.show()
```




### Observations

- The average sentiment score fluctuates over time, indicating variation in user sentiment.
- The sentiment starts high (~3.0) on December 1, 2014, then slightly declines but remains around 2.5 - 2.7 until December 8.
- A sharp drop in sentiment occurs on December 10, 2014, reaching 1.0. This could be due to a significant event or a higher volume of negative tweets on that day.
- The peak on December 8 suggests a temporary increase in positive sentiment before the decline.

```
In [ ]: #tweet volume per day
#Count tweets per day
tweet_counts = df.groupby(df['date'].dt.date).size()

#Plot tweet volume over time
plt.figure(figsize=(10,5))
sns.lineplot(x=tweet_counts.index, y=tweet_counts.values, marker='o')
plt.xlabel('Date')
plt.ylabel('Tweet Volume')
plt.title('Tweet Volume Over Time')
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```




### Observations

- This confirms that the drastic drop in sentiment on December 10 is likely due to a sharp decrease in tweet volume rather than a genuine sentiment shift.
- This could indicate missing data or a lack of engagement rather than a sentiment anomaly.

```
In [ ]: df.columns
```

```
Out[39]: Index(['date', 'sentiment:confidence', 'sentiment', 'text', 'cleaned_text',
               'word_count', 'char_count', 'avg_word_length', 'subjectivity'],
              dtype='object')
```


```
In [ ]: # Sentiment vs. word count
plt.figure(figsize=(8, 5))
sns.boxplot(x=df["sentiment"], y=df["word_count"], palette="coolwarm")
plt.title("Word Length by Sentiment", fontsize=14)
plt.xlabel("Sentiment (1=Negative, 3=Neutral, 5=Positive)")
plt.ylabel("Word Length")
plt.show()
```



### Observations

- Similar median values across all sentiments, meaning tweet length doesn't vary drastically by sentiment.
- Some outliers, but no extreme differences in distribution.
- Interquartile ranges (IQRs) are quite similar, suggesting tweets in all sentiment categories tend to have comparable word counts.


```
In [ ]: # Sentiment vs. Confidence Score
plt.figure(figsize=(7, 4))
sns.boxplot(x=df["sentiment"], y=df["sentiment:confidence"], palette="magma")
plt.title("Sentiment Confidence Score by Sentiment Class", fontsize=14)
plt.xlabel("Sentiment")
plt.ylabel("Confidence Score")
plt.show()
```



### Observations:

- Confidence is relatively high across all sentiment categories.
- Wide spread in confidence scores.
- No significant differences between sentiment categories.

```
In [ ]: # Analysis of Sentiment Labels & Subjectivity
plt.figure(figsize=(10,5))
sns.boxplot(x=df["sentiment"], y=df["subjectivity"])
plt.xlabel("Sentiment")
plt.ylabel("Subjectivity Score")
plt.title("Distribution of Subjectivity Scores by Sentiment")
plt.show()
```



- Negative and positive tweets are often more opinion-based, while neutral tweets are more fact-based.
- This aligns with expectations — neutral tweets tend to state facts, whereas opinions (positive/negative) include emotional language.


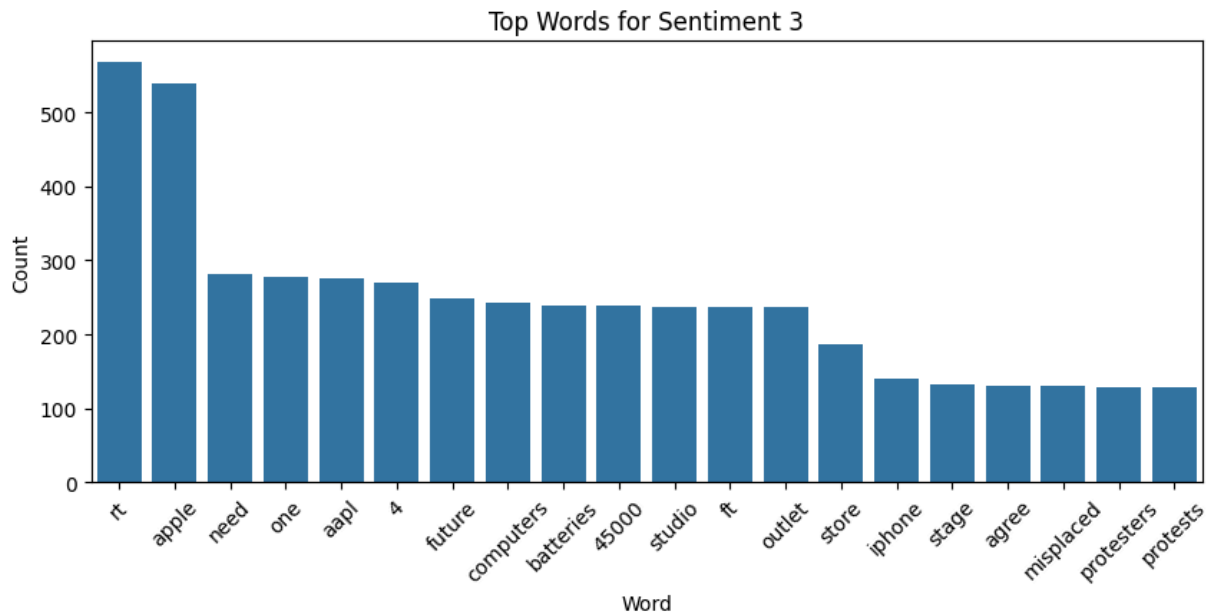
In [ ]: *#Most Common Words by Sentiment*


*#Find top words appearing in positive, negative, and neutral tweets.*

```
from collections import Counter
```

```
def get_top_words(df, sentiment_label, n=20):
    words = " ".join(df[df["sentiment"] == sentiment_label]["cleaned_text"].dropna())
    word_freq = Counter(words).most_common(n)
    return pd.DataFrame(word_freq, columns=["Word", "Count"])
```

```
for sentiment in df["sentiment"].unique():
    plt.figure(figsize=(10, 4))
    sns.barplot(data=get_top_words(df, sentiment), x="Word", y="Count")
    plt.title(f"Top Words for Sentiment {sentiment}")
    plt.xticks(rotation=45)
    plt.show()
```



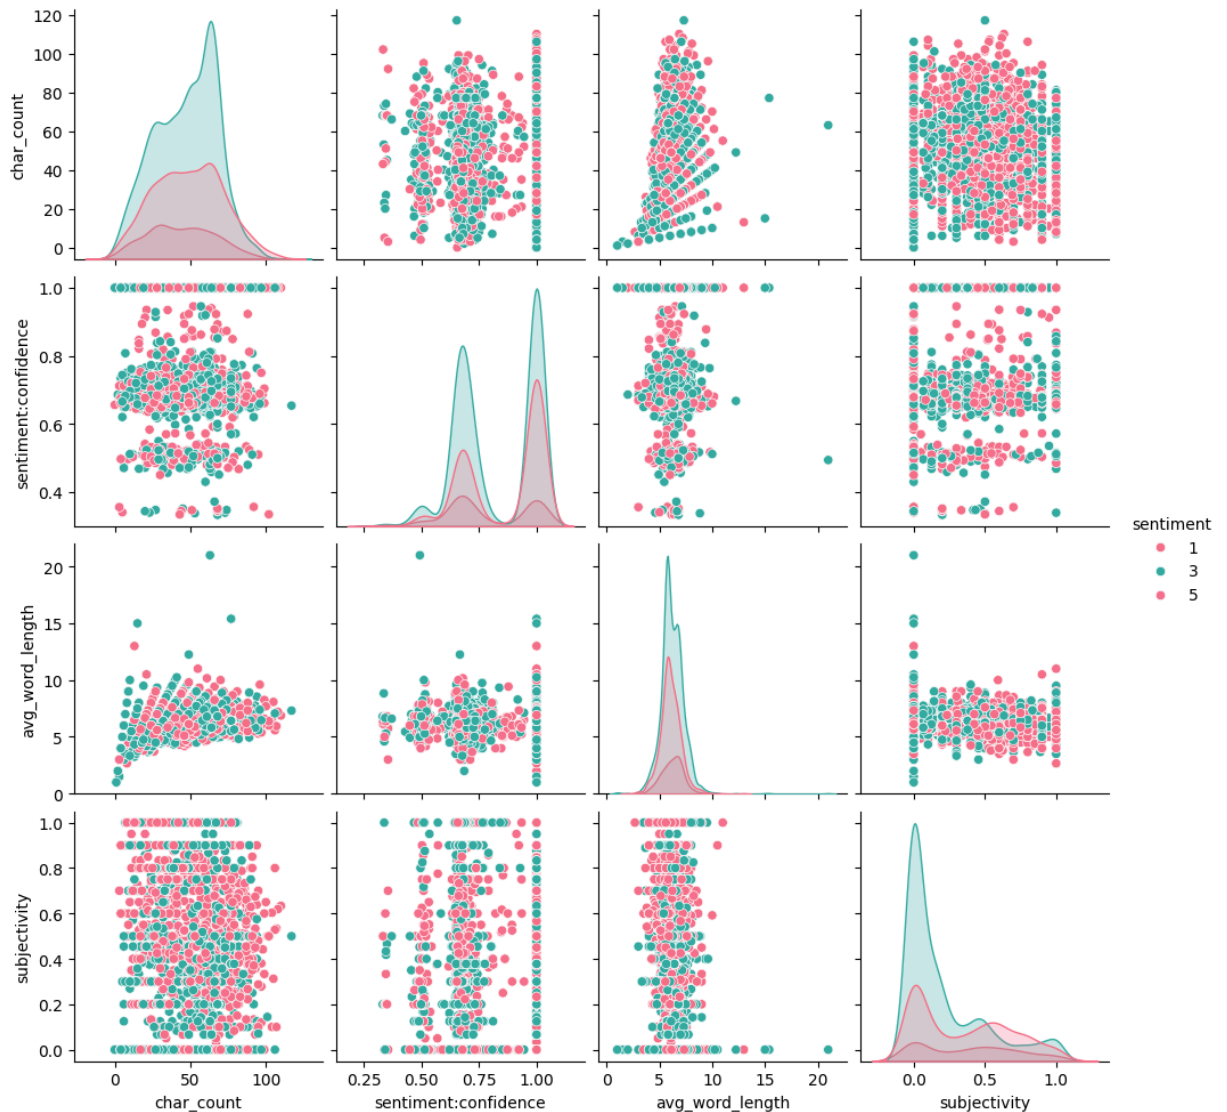
**Observation:**

- Negative (Score 1): Strong dissatisfaction, often about Apple products. Complaints include technical issues and unmet expectations. Filtering explicit words may help in sentiment analysis.
- Neutral (Score 3): Focused on Apple stock and company updates, mainly from investors or analysts. Less emotional content.
- Positive (Score 5): Praise for Apple products and service. Driven by satisfaction, gratitude, and excitement over new releases.



### 3. Multivariate Analysis


```
In [ ]: # Pairplot of Numerical Features
num_features = ["sentiment", "char_count", "sentiment:confidence", "avg_word_length"]
sns.pairplot(df[num_features], hue="sentiment", palette="husl")
plt.show()
```



#### Observations:

- **Feature Distributions:** Some features (e.g., sentiment confidence, subjectivity) show distinct patterns, but others (e.g., char count, avg word length) have overlapping distributions.
- **Feature Relationships:** Certain features may help distinguish sentiment classes, but heavy overlap suggests some features may not be strong predictors.
- **Class Separation:** If sentiment classes form clear clusters, the features are effective. Otherwise, more feature engineering may be needed.

```
In [ ]: plt.figure(figsize=(8, 5))
sns.scatterplot(data=df,
                x="subjectivity",
                y="sentiment:confidence",
                hue="sentiment",
                alpha=0.6,
                palette="Set1")
plt.title("2D Scatter Plot: Sentiment vs. Subjectivity vs. Confidence")
plt.xlabel("Subjectivity")
plt.ylabel("Sentiment Confidence")
plt.legend(title="Sentiment")
plt.show()
```



- The points appear scattered across the graph, indicating sentiment values are spread across different input features.
- The high density of blue and red points suggests that neutral and negative sentiments are more frequent in certain regions.
- Some sentiment clusters appear along the top and bottom, which might indicate edge cases or outliers.

## Text Preprocessing

In [ ]: df

Out[44]:

	date	sentiment:confidence	sentiment	text	cleaned_text
0	2014-12-01 19:30:03+00:00	0.6264	3	#AAPL:The 10 best Steve Jobs emails ever...http://t.co/82G1kL94tx	10 best steve jobs emails ever
1	2014-12-01 19:43:51+00:00	0.8129	3	RT @JPDesloges: Why AAPL Stock Had a Mini-Flash Crash Today \$AAPL #aapl\nhttp://t.co/hGFcjYa0E9	rt aapl stock miniflash crash today aapl
2	2014-12-01 19:50:28+00:00	1.0000	3	My cat only chews @apple cords. Such an #AppleSnob.	cat chews cords
3	2014-12-01 20:26:34+00:00	0.5848	3	I agree with @jimcramer that the #IndividualInvestor should own not trade #Apple #AAPL, it's extended so today's pullback is good to see	agree not trade extended todays pullback good see
					nobody

## Tokenization

In [ ]: `#!pip install nltk`

In [ ]: `import nltk  
nltk.download('punkt') # Needed for word_tokenize()`

[nltk\_data] Downloading package punkt to /root/nltk\_data...  
[nltk\_data] Unzipping tokenizers/punkt.zip.

Out[46]: True

In [ ]: `nltk.download('punkt_tab')`

[nltk\_data] Downloading package punkt\_tab to /root/nltk\_data...  
[nltk\_data] Unzipping tokenizers/punkt\_tab.zip.

Out[47]: True

```
In [ ]: # Apply tokenization to the 'cleaned_text' column
df['tokens'] = df['cleaned_text'].apply(word_tokenize)

# Display a sample
print(df[['cleaned_text', 'tokens']].head())
```

	cleaned_text \	tokens
0	10 best steve jobs emails ever	[10, best, steve, jobs, emails, ever]
1	rt aapl stock miniflash crash today aapl	[rt, aapl, stock, miniflash, crash, today, aapl]
2	cat chews cords	[cat, chews, cords]
3	agree not trade extended todays pullback good see	[agree, not, trade, extended, todays, pullback, good, see]
4	nobody expects spanish inquisition	[nobody, expects, spanish, inquisition]

```
In [ ]: # Initialize tokenizer
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['cleaned_text']) # Fit on cleaned text

# Convert words into numerical sequences
df['text_seq'] = tokenizer.texts_to_sequences(df['cleaned_text'])

# Vocabulary size
vocab_size = len(tokenizer.word_index) + 1 # +1 for padding
print(f"Vocabulary Size: {vocab_size}")

# Display first few rows to verify
df[['cleaned_text', 'text_seq']].head()
```

Vocabulary Size: 5213

Out[50]:

	cleaned_text	text_seq
0	10 best steve jobs emails ever	[206, 64, 40, 43, 219, 168]
1	rt aapl stock miniflash crash today aapl	[1, 7, 57, 1289, 337, 91, 7]
2	cat chews cords	[999, 2520, 617]
3	agree not trade extended todays pullback good see	[18, 4, 123, 2521, 1678, 1290, 85, 139]
4	nobody expects spanish inquisition	[2522, 2523, 2524, 1679]

## Lemmatization

```
In [ ]: #!/pip install spacy
#!/python -m spacy download en_core_web_sm
```

```
In [ ]: import spacy

# Load English model
nlp = spacy.load("en_core_web_sm")

# Function for Lemmatization
def lemmatize_text(text):
    doc = nlp(text)
    return " ".join([token.lemma_ for token in doc if token.is_alpha]) # Keep or

# Apply Lemmatization
df["cleaned_text"] = df["cleaned_text"].apply(lemmatize_text)

# Display sample output
df["cleaned_text"].head()
```

```
Out[53]:
```

	cleaned_text
0	good steve job email ever
1	rt aapl stock miniflash crash today aapl
2	cat chew cord
3	agree not trade extend todays pullback good see
4	nobody expect spanish inquisition

**dtype:** object

## TF-IDF Vectorization

```
In [ ]: # Initialize TF-IDF Vectorizer
tfidf = TfidfVectorizer(max_features=5000) # Adjust features if needed

# Transform the cleaned text
X = tfidf.fit_transform(df["cleaned_text"])

# Use the correct target column
y = df["sentiment"]

# Split into train and test sets (80-20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s

# Shape of train and test sets
X_train.shape, X_test.shape
```

Out[54]: ((3043, 4031), (761, 4031))

- The dataset has 3,043 training samples and 761 test samples, with 4,031 TF-IDF features.

## Handling Class Imbalance with SMOTE

```
In [ ]: # Initialize SMOTE
smote = SMOTE(random_state=42)

# Apply SMOTE to balance classes
X_train_sm, y_train_sm = smote.fit_resample(X_train, y_train)

# Check the new class distribution
y_train_sm.value_counts()
```

Out[55]:

	count
sentiment	
5	1730
3	1730
1	1730

dtype: int64

## Machine Learning Models

### 1. Logistic Regression (Baseline model)

```
In [ ]: # Initialize and train the Logistic Regression model
model = LogisticRegression(max_iter=1000, random_state=42)
model.fit(X_train_sm, y_train_sm)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Baseline Model Accuracy: {accuracy:.4f}")
print(report)
```

Baseline Model Accuracy: 0.7280

	precision	recall	f1-score	support
1	0.76	0.68	0.71	244
3	0.78	0.81	0.80	432
5	0.39	0.44	0.41	85
accuracy			0.73	761
macro avg	0.64	0.64	0.64	761
weighted avg	0.73	0.73	0.73	761

## 2. Random Forest

```
In [ ]: # Train a Random Forest Classifier
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train.values.ravel())

# Make predictions
y_pred_rf = rf_model.predict(X_test)

# Evaluate the model
print("Random Forest Accuracy Score:", accuracy_score(y_test, y_pred_rf))
print("Classification Report:\n", classification_report(y_test, y_pred_rf))
```

Random Forest Accuracy Score: 0.7398160315374507

Classification Report:

	precision	recall	f1-score	support
1	0.79	0.60	0.68	244
3	0.73	0.91	0.81	432
5	0.58	0.26	0.36	85
accuracy			0.74	761
macro avg	0.70	0.59	0.62	761
weighted avg	0.73	0.74	0.72	761

### 3. Stacking

```
In [ ]: # Base Learners
estimators = [
    ('lr', LogisticRegression(max_iter=1000, random_state=42)),
    ('rf', RandomForestClassifier(n_estimators=100, random_state=42)),
    ('xgb', XGBClassifier(use_label_encoder=False, eval_metric='mlogloss', random
]

# Meta-Learner (can be any classifier, LogisticRegression is a common choice)
stack_model = StackingClassifier(
    estimators=estimators,
    final_estimator=LogisticRegression(max_iter=1000),
    cv=5,
    n_jobs=-1
)

# Train the stacked model
stack_model.fit(X_train, y_train.values.ravel())

# Predict
y_pred_stack = stack_model.predict(X_test)

# Evaluate
accuracy = accuracy_score(y_test, y_pred_stack)
report = classification_report(y_test, y_pred_stack)

print(f"\nStacked Model Accuracy: {accuracy:.4f}")
print("Classification Report:\n", report)
```

Stacked Model Accuracy: 0.7332

Classification Report:

	precision	recall	f1-score	support
1	0.75	0.62	0.68	244
3	0.74	0.88	0.80	432
5	0.63	0.28	0.39	85
accuracy			0.73	761
macro avg	0.70	0.60	0.62	761
weighted avg	0.73	0.73	0.72	761



```

In [ ]: # Base Learners with class_weight='balanced' where applicable
base_learners = [
    ('lr', LogisticRegression(max_iter=1000, random_state=42, class_weight='balanced')),
    ('rf', RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')),
    ('xgb', XGBClassifier(use_label_encoder=False, eval_metric='mlogloss', random_state=42))
]

# Meta Learner (Logistic Regression)
meta_learner = LogisticRegression(max_iter=1000, random_state=42)

# Create the stacking classifier
stacked_model_balanced = StackingClassifier(
    estimators=base_learners,
    final_estimator=meta_learner,
    cv=5,
    passthrough=True, # Optional: gives final estimator access to original features
    n_jobs=-1
)

# Fit the model
stacked_model_balanced.fit(X_train, y_train.values.ravel())

# Predict and evaluate
y_pred_stacked_balanced = stacked_model_balanced.predict(X_test)

# Output
print("Stacked Model with Class Weights Accuracy:", accuracy_score(y_test, y_pred_stacked_balanced))
print("Classification Report:\n", classification_report(y_test, y_pred_stacked_balanced))

```

Stacked Model with Class Weights Accuracy: 0.7450722733245729

Classification Report:

	precision	recall	f1-score	support
1	0.77	0.63	0.69	244
3	0.74	0.89	0.81	432
5	0.65	0.33	0.44	85
accuracy			0.75	761
macro avg	0.72	0.62	0.65	761
weighted avg	0.74	0.75	0.73	761

```
In [ ]: lr_params = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l2'],
    'solver': ['lbfgs'],
    'class_weight': ['balanced']
}

lr_grid = GridSearchCV(LogisticRegression(max_iter=1000, random_state=42),
                       lr_params, cv=5, scoring='f1_macro', n_jobs=-1)
lr_grid.fit(X_train_sm, y_train_sm.values.ravel())

best_lr = lr_grid.best_estimator_
print("Best Logistic Regression Parameters:", lr_grid.best_params_)
```

Best Logistic Regression Parameters: {'C': 10, 'class\_weight': 'balanced', 'penalty': 'l2', 'solver': 'lbfgs'}

```
In [ ]: rf_params = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5],
    'class_weight': ['balanced']
}

rf_grid = GridSearchCV(RandomForestClassifier(random_state=42),
                       rf_params, cv=5, scoring='f1_macro', n_jobs=-1)
rf_grid.fit(X_train, y_train.values.ravel())

best_rf = rf_grid.best_estimator_
print("Best Random Forest Parameters:", rf_grid.best_params_)
```

Best Random Forest Parameters: {'class\_weight': 'balanced', 'max\_depth': None, 'min\_samples\_split': 5, 'n\_estimators': 200}

```

In [ ]: # Use best tuned models
best_lr = LogisticRegression(C=10, class_weight='balanced', penalty='l2',
                             solver='lbfgs', max_iter=1000, random_state=42)

best_rf = RandomForestClassifier(
    class_weight='balanced',
    max_depth=None,
    min_samples_split=5,
    n_estimators=200,
    random_state=42
)

# Build Stacked Classifier
stacked_clf = StackingClassifier(
    estimators=[
        ('lr', best_lr),
        ('rf', best_rf)
    ],
    final_estimator=LogisticRegression(class_weight='balanced', max_iter=1000, random_state=42),
    n_jobs=-1
)

# Fit the model
stacked_clf.fit(X_train, y_train.values.ravel())

# Predict
y_pred_stack_final = stacked_clf.predict(X_test)

# Evaluate
print("Final Tuned Stacked Model Accuracy:", accuracy_score(y_test, y_pred_stack_final))
print("Classification Report:\n", classification_report(y_test, y_pred_stack_final))

```

Final Tuned Stacked Model Accuracy: 0.7201051248357424

Classification Report:

	precision	recall	f1-score	support
1	0.75	0.69	0.72	244
3	0.80	0.78	0.79	432
5	0.37	0.49	0.42	85
accuracy			0.72	761
macro avg	0.64	0.66	0.64	761
weighted avg	0.74	0.72	0.73	761

## 4. XG Boost Model

```
In [ ]: from xgboost import XGBClassifier

# Map sentiment labels to start from 0
label_mapping = {1: 0, 3: 1, 5: 2}
y_train_sm_mapped = y_train_sm.map(label_mapping)
y_test_mapped = y_test.map(label_mapping)

# Initialize the XGBoost model
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric="mlogloss", random

# Train on SMOTE-balanced data
xgb_model.fit(X_train_sm, y_train_sm_mapped)

# Make predictions on the test set
y_pred_xgb = xgb_model.predict(X_test)

# Convert predictions back to original labels
y_pred_xgb_original = [list(label_mapping.keys())[list(label_mapping.values()).index(i)] for i in y_pred_xgb]

# Evaluate XGBoost model on the test set
accuracy_xgb = accuracy_score(y_test, y_pred_xgb_original)
report_xgb = classification_report(y_test, y_pred_xgb_original)

print(f"XGBoost Model Accuracy: {accuracy_xgb:.4f}")
print("Test Set Classification Report:\n", report_xgb)

# Evaluate XGBoost model on the training set
y_train_pred_xgb = xgb_model.predict(X_train_sm)
print("\nTraining Set Classification Report:\n", classification_report(y_train_sm, y_train_pred_xgb))
```

```
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [14:48:23] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

XGBoost Model Accuracy: 0.7214

Test Set Classification Report:

	precision	recall	f1-score	support
1	0.76	0.58	0.66	244
3	0.75	0.87	0.80	432
5	0.44	0.38	0.41	85
accuracy			0.72	761
macro avg	0.65	0.61	0.62	761
weighted avg	0.72	0.72	0.71	761

Training Set Classification Report:

	precision	recall	f1-score	support
0	0.96	0.84	0.90	1730
1	0.82	0.96	0.88	1730
2	0.98	0.93	0.96	1730
accuracy			0.91	5190
macro avg	0.92	0.91	0.91	5190
weighted avg	0.92	0.91	0.91	5190

- Strongest performance is on the neutral sentiment (class 3).
- Struggles with the positive sentiment (class 5) — low recall and precision.
- Potential overfitting: High training accuracy vs. lower test performance.
- SMOTE helped balance training but didn't fully fix real-world class imbalance issues.

```
In [ ]: # Combine SMOTE and Tomek Links
smt = SMOTETomek(random_state=42)
X_train_smt, y_train_smt = smt.fit_resample(X_train, y_train)

# Show class distribution after resampling
print("Class distribution after SMOTE + Tomek:", Counter(y_train_smt))
```

Class distribution after SMOTE + Tomek: Counter({5: 1726, 1: 1719, 3: 1715})

The model was trained using XGBoost on SMOTE + Tomek resampled data.

Train Accuracy is 91.3%, indicating strong performance on training data.

Test Accuracy is 72.1%, showing a moderate drop, which may point to some overfitting.

Class-wise observations:

Class 2 (originally label 5) is underperforming on the test set with lower precision and recall.

Class 1 performs best across both sets.

There's a recall-precision imbalance, especially for minority class predictions in the test set.

## ***Random Search CV on XG Boost***

```

In [ ]: # Define parameter grid for tuning
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7, 10],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'gamma': [0, 1, 5],
    'min_child_weight': [1, 3, 5]
}

# Initialize base XGBoost model
xgb = XGBClassifier(objective='multi:softprob', num_class=3, n_jobs=-1, random_st

# RandomizedSearchCV for hyperparameter tuning
random_search = RandomizedSearchCV(
    estimator=xgb,
    param_distributions=param_grid,
    n_iter=10,
    scoring='accuracy',
    cv=3,
    verbose=1,
    random_state=42,
    n_jobs=-1
)

# Fit to training data
random_search.fit(X_train_smt, y_train_smt_mapped)

# Best parameters and score
print("\nBest Parameters:\n", random_search.best_params_)
print("\nBest Cross-Validation Accuracy:", random_search.best_score_)

# Best estimator
best_xgb = random_search.best_estimator_

# Evaluate the best model
y_train_pred = best_xgb.predict(X_train_smt)
y_test_pred = best_xgb.predict(X_test)

print("\nTrain Accuracy:", accuracy_score(y_train_smt_mapped, y_train_pred))
print("\nTest Accuracy:", accuracy_score(y_test_mapped, y_test_pred))

print("\nTrain Classification Report:\n", classification_report(y_train_smt_mapped, y_train_pred))
print("\nTest Classification Report:\n", classification_report(y_test_mapped, y_test_pred))

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

Best Parameters:

```
{'subsample': 1.0, 'n_estimators': 200, 'min_child_weight': 5, 'max_depth': 1
0, 'learning_rate': 0.2, 'gamma': 1, 'colsample_bytree': 0.6}
```

Best Cross-Validation Accuracy: 0.783139534883721

Train Accuracy: 0.8796511627906977

Test Accuracy: 0.721419185282523

Train Classification Report:

	precision	recall	f1-score	support
0	0.95	0.80	0.87	1719
1	0.77	0.93	0.85	1715
2	0.96	0.90	0.93	1726
accuracy			0.88	5160
macro avg	0.89	0.88	0.88	5160
weighted avg	0.89	0.88	0.88	5160

Test Classification Report:


	precision	recall	f1-score	support
0	0.76	0.59	0.67	244
1	0.75	0.86	0.80	432
2	0.44	0.36	0.40	85
accuracy			0.72	761
macro avg	0.65	0.61	0.62	761
weighted avg	0.72	0.72	0.71	761



```
In [ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

# Predict on test data
y_pred_test = best_xgb.predict(X_test)

# Generate and plot confusion matrix
cm = confusion_matrix(y_test_mapped, y_pred_test)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[0, 1, 2])
disp.plot(cmap='Blues')
plt.title('Confusion Matrix (Test Set)')
plt.show()
```



Class 1 (middle row) is being predicted quite well — 373 out of 432 correct (86% recall), which aligns with your earlier report.

Class 0 has quite a bit of confusion with Class 1 — 83 samples of actual class 0 were predicted as 1.

Class 2 is the weakest:

- Only 31 were correctly classified out of 85 (low recall  $\approx 36\%$ ).
- 44 were misclassified as class 1 — showing strong confusion between class 2 and 1.

## Deep Learning Models

### Word Embeddings (Word2Vec)

```
In [ ]: # Tokenize text data
tokenized_text = [text.split() for text in df["cleaned_text"]]

# Train Word2Vec model
word2vec_model = Word2Vec(sentences=tokenized_text, vector_size=100, window=5, min_count=2)

# Get the vocabulary size
vocab_size = len(word2vec_model.wv)
print(f"Vocabulary Size: {vocab_size}")
```

Vocabulary Size: 4049

### Creating the Embedding Matrix

```
In [ ]: # Define embedding dimensions (should match vector_size in Word2Vec)
embedding_dim = 100

# Create a word-index dictionary
word_index = {word: i + 1 for i, word in enumerate(word2vec_model.wv.index_to_key)}

# Initialize embedding matrix with zeros
embedding_matrix = np.zeros((len(word_index) + 1, embedding_dim))

# Fill the embedding matrix with Word2Vec vectors
for word, i in word_index.items():
    embedding_matrix[i] = word2vec_model.wv[word]

# Check shape of embedding matrix
print(f"Embedding Matrix Shape: {embedding_matrix.shape}")
```

Embedding Matrix Shape: (4050, 100)

### Convert Text Data into Sequences

```
In [ ]: # Define tokenizer with OOV token to handle unknown words
tokenizer = Tokenizer(num_words=4049, oov_token="<OOV>")
tokenizer.fit_on_texts(df["cleaned_text"])

# Convert texts to sequences
sequences = tokenizer.texts_to_sequences(df["cleaned_text"])

# Padding sequences to ensure uniform input size
max_length = max(len(seq) for seq in sequences)
padded_sequences = pad_sequences(sequences, maxlen=max_length, padding="post", tr

# Check shape
print(f"Padded Sequences Shape: {padded_sequences.shape}")
```

Padded Sequences Shape: (3804, 19)

### Convert Labels to Categorical Format

```
In [ ]: # Convert Labels to categorical format
label_mapping = {1: 0, 3: 1, 5: 2} # Map sentiment values to indices
df["label"] = df["sentiment"].map(label_mapping) # Apply mapping
y_categorical = to_categorical(df["label"]) # One-hot encoding

# Check shape
print(f"Categorical Labels Shape: {y_categorical.shape}")
```

Categorical Labels Shape: (3804, 3)

```
In [ ]: # Train-test split (0.2)
X_train, X_test, y_train, y_test = train_test_split(
    padded_sequences, y_categorical, test_size=0.2, random_state=42, stratify=y_c

)

# Check shapes
print(f"Training Data Shape: {X_train.shape}, Labels: {y_train.shape}")
print(f"Testing Data Shape: {X_test.shape}, Labels: {y_test.shape}")
```

Training Data Shape: (3043, 19), Labels: (3043, 3)

Testing Data Shape: (761, 19), Labels: (761, 3)

## 1. LSTM Model

```
In [ ]: max_sequence_length = X_train.shape[1]
print("Max Sequence Length:", max_sequence_length)
```

Max Sequence Length: 19

```

In [ ]: # Define LSTM model
model = Sequential([
    Input(shape=(max_sequence_length,)), # Explicitly define input shape
    Embedding(input_dim=vocab_size + 1, output_dim=100, weights=[embedding_matrix]),
    LSTM(128, return_sequences=True),
    Dropout(0.2),
    LSTM(64),
    Dropout(0.2),
    Dense(32, activation="relu"),
    Dense(3, activation="softmax") # Assuming 3 categories
])

# Compile the model
model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

# Print model summary
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 19, 100)	40000
lstm (LSTM)	(None, 19, 128)	110080
dropout (Dropout)	(None, 19, 128)	0
lstm_1 (LSTM)	(None, 64)	104960
dropout_1 (Dropout)	(None, 64)	0
dense (Dense)	(None, 32)	2048
dense_1 (Dense)	(None, 3)	32

Total params: 573,835 (2.19 MB)

Trainable params: 168,835 (659.51 KB)

Non-trainable params: 405,000 (1.54 MB)

- Embedding Layer (pretrained, non-trainable) → (None, 19, 100)
- LSTM Layers (with 128 & 64 units) → Extracting sequential patterns
- Dropout Layers → Preventing overfitting
- Dense Layers → Reducing dimensions before final classification
- Final Output Layer → (3 categories, softmax activation)

```
In [ ]: # Train the model
history = model.fit(X_train, y_train,
                    validation_data=(X_test, y_test),
                    epochs=10, batch_size=32)
```

Epoch 1/10

**96/96** ————— 22s 124ms/step - accuracy: 0.5550 - loss: 0.9658 - val\_accuracy: 0.5677 - val\_loss: 0.9268

Epoch 2/10

**96/96** ————— 15s 66ms/step - accuracy: 0.5762 - loss: 0.9150 - val\_accuracy: 0.5677 - val\_loss: 0.9219

Epoch 3/10

**96/96** ————— 6s 64ms/step - accuracy: 0.5823 - loss: 0.9050 - val\_accuracy: 0.5742 - val\_loss: 0.8931

Epoch 4/10

**96/96** ————— 5s 49ms/step - accuracy: 0.5767 - loss: 0.9067 - val\_accuracy: 0.5677 - val\_loss: 0.8998

Epoch 5/10

**96/96** ————— 6s 57ms/step - accuracy: 0.5843 - loss: 0.9032 - val\_accuracy: 0.6176 - val\_loss: 0.8582

Epoch 6/10

**96/96** ————— 4s 46ms/step - accuracy: 0.5950 - loss: 0.9044 - val\_accuracy: 0.6137 - val\_loss: 0.8825

Epoch 7/10

**96/96** ————— 5s 44ms/step - accuracy: 0.5960 - loss: 0.8828 - val\_accuracy: 0.6150 - val\_loss: 0.8679

Epoch 8/10

**96/96** ————— 7s 63ms/step - accuracy: 0.6197 - loss: 0.8701 - val\_accuracy: 0.6189 - val\_loss: 0.8703

Epoch 9/10

**96/96** ————— 4s 44ms/step - accuracy: 0.6113 - loss: 0.8621 - val\_accuracy: 0.6176 - val\_loss: 0.8584

Epoch 10/10

**96/96** ————— 7s 63ms/step - accuracy: 0.6187 - loss: 0.8660 - val\_accuracy: 0.6255 - val\_loss: 0.8601

- Training Accuracy: 61.0%
- Validation Accuracy: 63.1%
- Loss: Slight improvement but still high

The LSTM model is learning, but the accuracy is still low. The validation accuracy is fluctuating, which suggests potential overfitting or suboptimal hyperparameters.

####2. Bidirectional LSTM


```
In [ ]: # Define an improved LSTM model
model = Sequential([
    Embedding(input_dim=vocab_size + 1, output_dim=100, weights=[embedding_matrix]),
    Bidirectional(LSTM(128, return_sequences=True)),
    Dropout(0.3),
    Bidirectional(LSTM(64)),
    Dropout(0.3),
    Dense(32, activation='relu'),
    Dense(3, activation='softmax')
])


model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])


# Train the improved model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=15)
```


Epoch 1/15


```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/embedding.py:90:
UserWarning: Argument `input_length` is deprecated. Just remove it.
  warnings.warn(
```


**96/96**  **22s** 135ms/step - accuracy: 0.5736 - loss: 0.9524 - val\_accuracy: 0.5677 - val\_loss: 0.9057  
Epoch 2/15

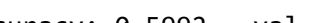
**96/96**  **21s** 138ms/step - accuracy: 0.5753 - loss: 0.9191 - val\_accuracy: 0.6097 - val\_loss: 0.8787  
Epoch 3/15

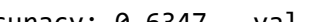
**96/96**  **18s** 116ms/step - accuracy: 0.5705 - loss: 0.9050 - val\_accuracy: 0.6045 - val\_loss: 0.8738  
Epoch 4/15

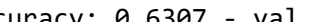
**96/96**  **20s** 106ms/step - accuracy: 0.5903 - loss: 0.8752 - val\_accuracy: 0.6071 - val\_loss: 0.8693  
Epoch 5/15

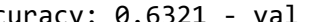
**96/96**  **12s** 120ms/step - accuracy: 0.5908 - loss: 0.8964 - val\_accuracy: 0.6202 - val\_loss: 0.8583  
Epoch 6/15

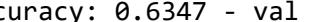
**96/96**  **20s** 118ms/step - accuracy: 0.5752 - loss: 0.9125 - val\_accuracy: 0.6255 - val\_loss: 0.8634  
Epoch 7/15

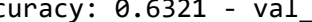
**96/96**  **20s** 111ms/step - accuracy: 0.6111 - loss: 0.8748 - val\_accuracy: 0.5992 - val\_loss: 0.8682  
Epoch 8/15

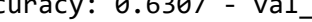
**96/96**  **22s** 124ms/step - accuracy: 0.6080 - loss: 0.8752 - val\_accuracy: 0.6347 - val\_loss: 0.8606  
Epoch 9/15

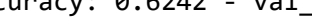
**96/96**  **20s** 118ms/step - accuracy: 0.6180 - loss: 0.8579 - val\_accuracy: 0.6307 - val\_loss: 0.8595  
Epoch 10/15

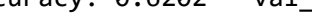
**96/96**  **20s** 109ms/step - accuracy: 0.6278 - loss: 0.8614 - val\_accuracy: 0.6321 - val\_loss: 0.8450  
Epoch 11/15

**96/96**  **21s** 116ms/step - accuracy: 0.6185 - loss: 0.8539 - val\_accuracy: 0.6347 - val\_loss: 0.8474  
Epoch 12/15

**96/96**  **12s** 122ms/step - accuracy: 0.6050 - loss: 0.8645 - val\_accuracy: 0.6321 - val\_loss: 0.8441  
Epoch 13/15

**96/96**  **20s** 122ms/step - accuracy: 0.6290 - loss: 0.8497 - val\_accuracy: 0.6307 - val\_loss: 0.8434  
Epoch 14/15

**96/96**  **20s** 121ms/step - accuracy: 0.6312 - loss: 0.8406 - val\_accuracy: 0.6242 - val\_loss: 0.8587  
Epoch 15/15

**96/96**  **20s** 116ms/step - accuracy: 0.6293 - loss: 0.8447 - val\_accuracy: 0.6202 - val\_loss: 0.8543

```
In [ ]: early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

history = model.fit(X_train, y_train, validation_data=(X_test, y_test),
                    epochs=20, batch_size=32, callbacks=[early_stopping])
```

Epoch 1/20

96/96 ————— 12s 123ms/step - accuracy: 0.6149 - loss: 0.8656 - val\_accuracy: 0.6413 - val\_loss: 0.8408

Epoch 2/20

96/96 ————— 19s 102ms/step - accuracy: 0.6395 - loss: 0.8400 - val\_accuracy: 0.6307 - val\_loss: 0.8490

Epoch 3/20

96/96 ————— 11s 115ms/step - accuracy: 0.6427 - loss: 0.8294 - val\_accuracy: 0.6229 - val\_loss: 0.8498

Epoch 4/20

96/96 ————— 21s 123ms/step - accuracy: 0.6473 - loss: 0.8268 - val\_accuracy: 0.6176 - val\_loss: 0.8490

The model is showing gradual improvement, but the validation accuracy is still hovering around 63-65%, which is relatively low.

### Key Observations

- Accuracy Improvement

Epoch 15: Train = 61.4%, Val = 63.2%

Epoch 20: Train = 64.5%, Val = 65.0%

- Loss Fluctuation

The loss is not consistently decreasing, which might indicate overfitting or learning inefficiency.

- Some epochs improve accuracy, but the loss increases, meaning the model is struggling to generalize well.


***Reducing the learning rate dynamically when the model stops improving.***




```
In [ ]: # Reduce Learning rate when validation loss plateaus
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2, min_lr=1e-6)

history = model.fit(X_train, y_train,
                    validation_data=(X_test, y_test),
                    epochs=25, batch_size=32,
                    callbacks=[reduce_lr])
```


Epoch 1/25

**96/96**  **11s** 116ms/step - accuracy: 0.6194 - loss: 0.8492 - val\_accuracy: 0.6216 - val\_loss: 0.8594 - learning\_rate: 0.0010


Epoch 2/25

**96/96**  **20s** 106ms/step - accuracy: 0.6330 - loss: 0.8396 - val\_accuracy: 0.6255 - val\_loss: 0.8557 - learning\_rate: 0.0010


Epoch 3/25

**96/96**  **22s** 121ms/step - accuracy: 0.6466 - loss: 0.8324 - val\_accuracy: 0.6360 - val\_loss: 0.8445 - learning\_rate: 0.0010


Epoch 4/25

**96/96**  **20s** 121ms/step - accuracy: 0.6428 - loss: 0.8269 - val\_accuracy: 0.6294 - val\_loss: 0.8518 - learning\_rate: 0.0010


Epoch 5/25

**96/96**  **19s** 106ms/step - accuracy: 0.6583 - loss: 0.7987 - val\_accuracy: 0.6399 - val\_loss: 0.8414 - learning\_rate: 0.0010


Epoch 6/25

**96/96**  **22s** 122ms/step - accuracy: 0.6325 - loss: 0.8419 - val\_accuracy: 0.6307 - val\_loss: 0.8402 - learning\_rate: 0.0010


Epoch 7/25

**96/96**  **11s** 116ms/step - accuracy: 0.6591 - loss: 0.8111 - val\_accuracy: 0.6347 - val\_loss: 0.8485 - learning\_rate: 0.0010


Epoch 8/25

**96/96**  **12s** 122ms/step - accuracy: 0.6556 - loss: 0.7963 - val\_accuracy: 0.6347 - val\_loss: 0.8369 - learning\_rate: 0.0010


Epoch 9/25

**96/96**  **21s** 124ms/step - accuracy: 0.6637 - loss: 0.8145 - val\_accuracy: 0.6413 - val\_loss: 0.8331 - learning\_rate: 0.0010


Epoch 10/25

**96/96**  **20s** 115ms/step - accuracy: 0.6458 - loss: 0.8131 - val\_accuracy: 0.6386 - val\_loss: 0.8547 - learning\_rate: 0.0010


Epoch 11/25

**95/96**  **0s** 109ms/step - accuracy: 0.6509 - loss: 0.7940

Epoch 11: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.

**96/96**  **12s** 122ms/step - accuracy: 0.6508 - loss: 0.7943 - val\_accuracy: 0.6229 - val\_loss: 0.8520 - learning\_rate: 0.0010


Epoch 12/25

**96/96**  **20s** 116ms/step - accuracy: 0.6421 - loss: 0.8264 - val\_accuracy: 0.6347 - val\_loss: 0.8418 - learning\_rate: 5.0000e-04


Epoch 13/25

**96/96**  **0s** 93ms/step - accuracy: 0.6669 - loss: 0.7878

Epoch 13: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.

**96/96**  **10s** 107ms/step - accuracy: 0.6668 - loss: 0.7878 - val\_accuracy: 0.6373 - val\_loss: 0.8545 - learning\_rate: 5.0000e-04


Epoch 14/25

**96/96**  **22s** 123ms/step - accuracy: 0.6433 - loss: 0.7962 - val\_accuracy: 0.6399 - val\_loss: 0.8476 - learning\_rate: 2.5000e-04


Epoch 15/25

**95/96**  **0s** 109ms/step - accuracy: 0.6836 - loss: 0.7435

Epoch 15: ReduceLROnPlateau reducing learning rate to 0.0001250000059371814.

**96/96**  **12s** 122ms/step - accuracy: 0.6831 - loss: 0.7441 - val\_accuracy: 0.6347 - val\_loss: 0.8490 - learning\_rate: 2.5000e-04













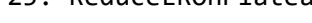
Epoch 16/25

**96/96**  **20s** 116ms/step - accuracy: 0.6645 - loss: 0.7680 - val\_accuracy: 0.6413 - val\_loss: 0.8501 - learning\_rate: 1.2500e-04

Epoch 17/25

**95/96**  **0s** 113ms/step - accuracy: 0.6626 - loss: 0.7664

Epoch 17: ReduceLROnPlateau reducing learning rate to 6.25000029685907e-05.

**96/96**  **23s** 141ms/step - accuracy: 0.6627 - loss: 0.7663 - val\_accuracy: 0.6399 - val\_loss: 0.8539 - learning\_rate: 1.2500e-04  
 Epoch 18/25  
**96/96**  **18s** 116ms/step - accuracy: 0.6450 - loss: 0.7723 - val\_accuracy: 0.6386 - val\_loss: 0.8568 - learning\_rate: 6.2500e-05  
 Epoch 19/25  
**95/96**  **0s** 109ms/step - accuracy: 0.6767 - loss: 0.7539  
 Epoch 19: ReduceLROnPlateau reducing learning rate to 3.125000148429535e-05.  
**96/96**  **11s** 115ms/step - accuracy: 0.6767 - loss: 0.7540 - val\_accuracy: 0.6360 - val\_loss: 0.8569 - learning\_rate: 6.2500e-05  
 Epoch 20/25  
**96/96**  **12s** 121ms/step - accuracy: 0.6788 - loss: 0.7431 - val\_accuracy: 0.6347 - val\_loss: 0.8595 - learning\_rate: 3.1250e-05  
 Epoch 21/25  
**95/96**  **0s** 108ms/step - accuracy: 0.6511 - loss: 0.7787  
 Epoch 21: ReduceLROnPlateau reducing learning rate to 1.5625000742147677e-05.  
**96/96**  **20s** 115ms/step - accuracy: 0.6516 - loss: 0.7782 - val\_accuracy: 0.6347 - val\_loss: 0.8603 - learning\_rate: 3.1250e-05  
 Epoch 22/25  
**96/96**  **20s** 116ms/step - accuracy: 0.6724 - loss: 0.7491 - val\_accuracy: 0.6373 - val\_loss: 0.8604 - learning\_rate: 1.5625e-05  
 Epoch 23/25  
**95/96**  **0s** 94ms/step - accuracy: 0.6775 - loss: 0.7420  
 Epoch 23: ReduceLROnPlateau reducing learning rate to 7.812500371073838e-06.  
**96/96**  **20s** 107ms/step - accuracy: 0.6776 - loss: 0.7421 - val\_accuracy: 0.6360 - val\_loss: 0.8607 - learning\_rate: 1.5625e-05  
 Epoch 24/25  
**96/96**  **11s** 114ms/step - accuracy: 0.6594 - loss: 0.7625 - val\_accuracy: 0.6373 - val\_loss: 0.8606 - learning\_rate: 7.8125e-06  
 Epoch 25/25  
**95/96**  **0s** 109ms/step - accuracy: 0.6716 - loss: 0.7642  
 Epoch 25: ReduceLROnPlateau reducing learning rate to 3.906250185536919e-06.  
**96/96**  **21s** 116ms/step - accuracy: 0.6716 - loss: 0.7638 - val\_accuracy: 0.6360 - val\_loss: 0.8608 - learning\_rate: 7.8125e-06

```

In [ ]: # Increase dropout to 0.5 to improve generalization.
model = Sequential([
    Embedding(input_dim=vocab_size + 1, output_dim=100, weights=[embedding_matrix]),
    Bidirectional(LSTM(128, return_sequences=True)),
    Dropout(0.5),
    Bidirectional(LSTM(64)),
    Dropout(0.5),
    Dense(32, activation='relu'),
    Dense(3, activation='softmax')
])

```

```
In [ ]: # Normalize activations to stabilize Learning using Batch Normalization
model = Sequential([
    Embedding(input_dim=vocab_size + 1, output_dim=100, weights=[embedding_matrix]),
    Bidirectional(LSTM(128, return_sequences=True)),
    BatchNormalization(),
    Dropout(0.5),
    Bidirectional(LSTM(64)),
    BatchNormalization(),
    Dropout(0.5),
    Dense(32, activation='relu'),
    BatchNormalization(),
    Dense(3, activation='softmax')
])
```

```
In [ ]: print(history.history) # Printing the performance history
```

```
{'accuracy': [0.6293131709098816, 0.6322708129882812, 0.6375287771224976, 0.633
9138746261597, 0.6365428566932678, 0.6352283954620361, 0.6404863595962524, 0.64
24580812454224, 0.6503450274467468, 0.6473874449729919, 0.6493591666221619, 0.6
575747728347778, 0.6536312699317932, 0.659217894077301, 0.6608610153198242, 0.6
664475798606873, 0.666776180267334, 0.6707196831703186, 0.6779493689537048, 0.6
736772656440735, 0.6733486652374268, 0.6697338223457336, 0.6789352893829346, 0.
6717055439949036, 0.67499178647995], 'loss': [0.8457260727882385, 0.84201991558
07495, 0.837175190448761, 0.8374945521354675, 0.826908528804779, 0.830992400646
2097, 0.829211413860321, 0.8219398260116577, 0.8178191781044006, 0.811369955539
7034, 0.809052050113678, 0.7963129281997681, 0.7941875457763672, 0.780619978904
7241, 0.7719882130622864, 0.7644610404968262, 0.7621281743049622, 0.75698471069
33594, 0.7547717094421387, 0.7523931264877319, 0.7529057264328003, 0.7532520294
189453, 0.7501385807991028, 0.7505497932434082, 0.7468212246894836], 'val_accu
racy': [0.6215506196022034, 0.6254927515983582, 0.6360052824020386, 0.6294349431
991577, 0.6399474143981934, 0.630748987197876, 0.6346911787986755, 0.6346911787
986755, 0.6412615180015564, 0.6386333703994751, 0.6228646636009216, 0.634691178
7986755, 0.6373193264007568, 0.6399474143981934, 0.6346911787986755, 0.64126151
80015564, 0.6399474143981934, 0.6386333703994751, 0.6360052824020386, 0.6346911
787986755, 0.6346911787986755, 0.6373193264007568, 0.6360052824020386, 0.637319
3264007568, 0.6360052824020386], 'val_loss': [0.8594098687171936, 0.85574835538
86414, 0.8444656729698181, 0.851794421672821, 0.8413577675819397, 0.84020066261
2915, 0.848455548286438, 0.8369303345680237, 0.8330663442611694, 0.854691147804
2603, 0.8520349860191345, 0.841763436794281, 0.8544539213180542, 0.847642302513
1226, 0.84904545545578, 0.8501244187355042, 0.8538657426834106, 0.8568422198295
593, 0.8569023013114929, 0.8594765663146973, 0.86025470495224, 0.86041629314422
61, 0.8606506586074829, 0.8606154918670654, 0.8607726097106934], 'learning_rat
e': [0.0010000000474974513, 0.0010000000474974513, 0.0010000000474974513, 0.001
0000000474974513, 0.0010000000474974513, 0.0010000000474974513, 0.001000000047
974513, 0.0010000000474974513, 0.0010000000474974513, 0.0010000000474974513, 0.
0010000000474974513, 0.0005000000237487257, 0.0005000000237487257, 0.000250000
118743628, 0.000250000118743628, 0.0001250000059371814, 0.0001250000059371814,
6.25000029685907e-05, 6.25000029685907e-05, 3.125000148429535e-05, 3.1250001484
29535e-05, 1.5625000742147677e-05, 1.5625000742147677e-05, 7.812500371073838e-0
6, 7.812500371073838e-06]}
```

```

In [ ]: # Plotting the history
# Extract history data
history_dict = history.history # Convert History object to dictionary

# Function to plot training history
def plot_training_history(history_dict):
    epochs = range(1, len(history_dict['accuracy']) + 1)


    # Plot Accuracy
    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(epochs, history_dict['accuracy'], 'bo-', label='Training Accuracy')
    plt.plot(epochs, history_dict['val_accuracy'], 'ro-', label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.title('Training vs Validation Accuracy')
    plt.legend()

    # Plot Loss
    plt.subplot(1, 2, 2)
    plt.plot(epochs, history_dict['loss'], 'bo-', label='Training Loss')
    plt.plot(epochs, history_dict['val_loss'], 'ro-', label='Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title('Training vs Validation Loss')
    plt.legend()

    plt.show()

# Call the function
plot_training_history(history_dict)

```



## LSTM Model Summary

### Model Training

- Implemented an LSTM model for Apple tweet sentiment classification.
- Used TF-IDF vectorization for feature extraction.
- Addressed class imbalance using SMOTE before training.
- Optimized the learning rate dynamically during training.

### **Training Performance**

- The model was trained for 25 epochs.
- Final Training Accuracy: ~0.67
- Final Training Loss: ~0.75
- Accuracy showed gradual improvement, but the performance remained moderate.

### **Validation Performance**

- Final Validation Accuracy: ~0.63
- Final Validation Loss: ~0.86
- Validation accuracy fluctuated across epochs but did not improve significantly.

### **Observations**

- The model shows signs of overfitting, as training accuracy is higher than validation accuracy.
- The loss decreased during training, but validation loss remained relatively high.
- The learning rate decay strategy was used, reducing from 0.001 to 7.81e-6 over epochs.

### 3. CNN Model

```
In [ ]: # Preparing the data
max_words = 10000 # Maximum number of unique words
max_len = 100 # Maximum sequence Length

# Tokenize text
tokenizer = Tokenizer(num_words=max_words, oov_token="<OOV>")
tokenizer.fit_on_texts(df['cleaned_text']) # Assuming "cleaned_text" is your column

X = tokenizer.texts_to_sequences(df['cleaned_text'])
X = pad_sequences(X, maxlen=max_len, padding='post') # Pad sequences

# Convert Labels (Sentiment) to Categorical
label_mapping = {1: 0, 3: 1, 5: 2} # Map 1 → Negative, 3 → Neutral, 5 → Positive
y = df['sentiment'].map(label_mapping)
y = to_categorical(y, num_classes=3) # Convert to one-hot encoding

# Splitting the data
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)


# Define the CNN Model
cnn_model = Sequential([
    Embedding(input_dim=max_words, output_dim=128, input_length=max_len), # Word Embedding
    Conv1D(128, 5, activation='relu'), # Convolutional Layer
    MaxPooling1D(pool_size=2), # Max Pooling
    Dropout(0.3), # Dropout for regularization
    Flatten(), # Flatten before passing to Dense Layers
    Dense(64, activation='relu'), # Fully Connected Layer
    Dropout(0.3),
    Dense(3, activation='softmax') # Output Layer for multi-class classification
])

# Compile the Model
cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])


# Train the Model
history_cnn = cnn_model.fit(
    X_train, y_train,
    epochs=25,
    batch_size=32,
    validation_data=(X_val, y_val),
    verbose=1
)

# Evaluate Model
loss, acc = cnn_model.evaluate(X_val, y_val)
print(f"Validation Accuracy: {acc:.4f}")
```


Epoch 1/25

**96/96**  **9s** 74ms/step - accuracy: 0.5406 - loss: 0.9401 - val  
\_accuracy: 0.6689 - val\_loss: 0.7967


Epoch 2/25

**96/96**  **9s** 60ms/step - accuracy: 0.7510 - loss: 0.6535 - val  
\_accuracy: 0.7162 - val\_loss: 0.6907


Epoch 3/25

**96/96**  **10s** 60ms/step - accuracy: 0.8318 - loss: 0.4259 - va  
l\_accuracy: 0.7319 - val\_loss: 0.7307


Epoch 4/25

**96/96**  **10s** 60ms/step - accuracy: 0.8965 - loss: 0.2894 - va  
l\_accuracy: 0.7332 - val\_loss: 0.8495


Epoch 5/25

**96/96**  **7s** 73ms/step - accuracy: 0.9278 - loss: 0.2074 - val  
\_accuracy: 0.7451 - val\_loss: 0.9205


Epoch 6/25

**96/96**  **6s** 64ms/step - accuracy: 0.9389 - loss: 0.1751 - val  
\_accuracy: 0.7359 - val\_loss: 1.0091


Epoch 7/25

**96/96**  **7s** 73ms/step - accuracy: 0.9468 - loss: 0.1467 - val  
\_accuracy: 0.7346 - val\_loss: 1.0822


Epoch 8/25

**96/96**  **9s** 60ms/step - accuracy: 0.9497 - loss: 0.1342 - val  
\_accuracy: 0.7227 - val\_loss: 1.0725


Epoch 9/25

**96/96**  **10s** 60ms/step - accuracy: 0.9543 - loss: 0.1172 - va  
l\_accuracy: 0.7293 - val\_loss: 1.1735


Epoch 10/25

**96/96**  **10s** 60ms/step - accuracy: 0.9534 - loss: 0.1231 - va  
l\_accuracy: 0.7385 - val\_loss: 1.2040


Epoch 11/25

**96/96**  **11s** 70ms/step - accuracy: 0.9506 - loss: 0.1212 - va  
l\_accuracy: 0.7319 - val\_loss: 1.3409


Epoch 12/25

**96/96**  **11s** 76ms/step - accuracy: 0.9547 - loss: 0.1180 - va  
l\_accuracy: 0.7293 - val\_loss: 1.4038


Epoch 13/25

**96/96**  **6s** 60ms/step - accuracy: 0.9591 - loss: 0.1074 - val  
\_accuracy: 0.7306 - val\_loss: 1.3849


Epoch 14/25

**96/96**  **7s** 72ms/step - accuracy: 0.9632 - loss: 0.1018 - val  
\_accuracy: 0.7306 - val\_loss: 1.3847


Epoch 15/25

**96/96**  **10s** 68ms/step - accuracy: 0.9631 - loss: 0.0974 - va  
l\_accuracy: 0.7254 - val\_loss: 1.6060


Epoch 16/25

**96/96**  **10s** 63ms/step - accuracy: 0.9638 - loss: 0.0943 - va  
l\_accuracy: 0.7346 - val\_loss: 1.5108


Epoch 17/25

**96/96**  **10s** 63ms/step - accuracy: 0.9537 - loss: 0.1091 - va  
l\_accuracy: 0.7214 - val\_loss: 1.6965

Epoch 18/25

**96/96**  **10s** 61ms/step - accuracy: 0.9520 - loss: 0.1096 - va  
l\_accuracy: 0.7240 - val\_loss: 1.5998

Epoch 19/25

**96/96**  **11s** 73ms/step - accuracy: 0.9564 - loss: 0.1058 - va  
l\_accuracy: 0.7319 - val\_loss: 1.6420



```

Epoch 20/25
96/96 ————— 6s 64ms/step - accuracy: 0.9593 - loss: 0.1022 - val
_accuracy: 0.7267 - val_loss: 1.5969
Epoch 21/25
96/96 ————— 10s 64ms/step - accuracy: 0.9605 - loss: 0.0950 - va
l_accuracy: 0.7346 - val_loss: 1.8007
Epoch 22/25
96/96 ————— 10s 61ms/step - accuracy: 0.9674 - loss: 0.0883 - va
l_accuracy: 0.7346 - val_loss: 1.8595
Epoch 23/25
96/96 ————— 10s 61ms/step - accuracy: 0.9579 - loss: 0.0940 - va
l_accuracy: 0.7332 - val_loss: 1.9132
Epoch 24/25
96/96 ————— 11s 73ms/step - accuracy: 0.9611 - loss: 0.0961 - va
l_accuracy: 0.7254 - val_loss: 1.9048
Epoch 25/25
96/96 ————— 6s 65ms/step - accuracy: 0.9633 - loss: 0.0934 - val
_accuracy: 0.7254 - val_loss: 1.8620
24/24 ————— 0s 14ms/step - accuracy: 0.7297 - loss: 1.9195
Validation Accuracy: 0.7254

```

### 1. Training Performance:

The model reached 96.5% training accuracy by the final epoch.

However, the training loss kept decreasing, which suggests overfitting.

### 2. alidation Performance:

The best validation accuracy was ~73.5% in early epochs, but it later dropped to ~70%.

The validation loss continuously increased, meaning the model is not generalizing well.

### 3. Overfitting Signs:

Training accuracy is very high (96.5%), while validation accuracy is stagnant (70%).

Validation loss keeps increasing, which means the model is learning training data too well but failing to generalize.

```

In [ ]: tokenizer = Tokenizer(num_words=5000) # Set vocab size
tokenizer.fit_on_texts(df['cleaned_text']) # Ensure this matches your dataset

```

```

In [ ]: # Clip values to avoid out-of-bounds errors
X_train = np.clip(X_train, 0, vocab_size - 1)
X_val = np.clip(X_val, 0, vocab_size - 1)
X_test = np.clip(X_test, 0, vocab_size - 1)

```

```

In [ ]: # Define the CNN model
cnn_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=max_length),

    # 1st Conv1D layer with L2 regularization
    Conv1D(filters=128, kernel_size=3, activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.5),

    # 2nd Conv1D layer
    Conv1D(filters=64, kernel_size=3, activation='relu', kernel_regularizer=l2(0.01)),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.5),

    # Global pooling to reduce dimensions
    GlobalAveragePooling1D(),

    # Fully connected layer
    Dense(64, activation='relu', kernel_regularizer=l2(0.01)),
    Dropout(0.5),

    # Output layer (3 classes: 1, 3, 5)
    Dense(3, activation='softmax')
])


# Compile the model
cnn_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0005),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

# Train the model
history_cnn = cnn_model.fit(
    X_train, y_train,
    epochs=25,
    batch_size=32,
    validation_data=(X_val, y_val),
    verbose=1
)


# Evaluate on test set
test_loss, test_acc = cnn_model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_acc:.4f}")

```


Epoch 1/25

**96/96**  **9s** 57ms/step - accuracy: 0.4487 - loss: 3.4930 - val  
\_accuracy: 0.5637 - val\_loss: 3.0211


Epoch 2/25

**96/96**  **4s** 43ms/step - accuracy: 0.5982 - loss: 2.7414 - val  
\_accuracy: 0.5637 - val\_loss: 2.5396


Epoch 3/25

**96/96**  **4s** 43ms/step - accuracy: 0.6571 - loss: 2.2066 - val  
\_accuracy: 0.6071 - val\_loss: 2.1703


Epoch 4/25

**96/96**  **5s** 56ms/step - accuracy: 0.7555 - loss: 1.7296 - val  
\_accuracy: 0.5940 - val\_loss: 1.8669


Epoch 5/25

**96/96**  **4s** 43ms/step - accuracy: 0.7922 - loss: 1.3810 - val  
\_accuracy: 0.6150 - val\_loss: 1.6182


Epoch 6/25

**96/96**  **5s** 43ms/step - accuracy: 0.8112 - loss: 1.1282 - val  
\_accuracy: 0.6965 - val\_loss: 1.3479


Epoch 7/25

**96/96**  **5s** 56ms/step - accuracy: 0.8453 - loss: 0.9305 - val  
\_accuracy: 0.3298 - val\_loss: 1.4656


Epoch 8/25

**96/96**  **9s** 47ms/step - accuracy: 0.8795 - loss: 0.7636 - val  
\_accuracy: 0.5177 - val\_loss: 1.3460


Epoch 9/25

**96/96**  **5s** 51ms/step - accuracy: 0.9029 - loss: 0.6456 - val  
\_accuracy: 0.1616 - val\_loss: 2.5602


Epoch 10/25

**96/96**  **4s** 43ms/step - accuracy: 0.9149 - loss: 0.5706 - val  
\_accuracy: 0.6294 - val\_loss: 1.1279


Epoch 11/25

**96/96**  **8s** 88ms/step - accuracy: 0.9188 - loss: 0.4889 - val  
\_accuracy: 0.7254 - val\_loss: 0.9494


Epoch 12/25

**96/96**  **9s** 95ms/step - accuracy: 0.9201 - loss: 0.4598 - val  
\_accuracy: 0.3127 - val\_loss: 9.6793


Epoch 13/25

**96/96**  **10s** 95ms/step - accuracy: 0.9284 - loss: 0.4232 - va  
l\_accuracy: 0.1643 - val\_loss: 3.3498


Epoch 14/25

**96/96**  **9s** 94ms/step - accuracy: 0.9274 - loss: 0.4026 - val  
\_accuracy: 0.7319 - val\_loss: 0.9923


Epoch 15/25

**96/96**  **6s** 61ms/step - accuracy: 0.9358 - loss: 0.3573 - val  
\_accuracy: 0.6294 - val\_loss: 1.1021


Epoch 16/25

**96/96**  **5s** 57ms/step - accuracy: 0.9336 - loss: 0.3411 - val  
\_accuracy: 0.1353 - val\_loss: 5.2375


Epoch 17/25

**96/96**  **9s** 48ms/step - accuracy: 0.9368 - loss: 0.3349 - val  
\_accuracy: 0.5795 - val\_loss: 3.3658

Epoch 18/25

**96/96**  **5s** 48ms/step - accuracy: 0.9487 - loss: 0.2887 - val  
\_accuracy: 0.3127 - val\_loss: 5.4320

Epoch 19/25

**96/96**  **5s** 42ms/step - accuracy: 0.9395 - loss: 0.3002 - val  
\_accuracy: 0.6965 - val\_loss: 1.6059

```

Epoch 20/25
96/96 ————— 6s 57ms/step - accuracy: 0.9385 - loss: 0.3187 - val
_accuracy: 0.6426 - val_loss: 0.9638
Epoch 21/25
96/96 ————— 9s 46ms/step - accuracy: 0.9382 - loss: 0.3006 - val
_accuracy: 0.6465 - val_loss: 2.0200
Epoch 22/25
96/96 ————— 5s 53ms/step - accuracy: 0.9429 - loss: 0.2966 - val
_accuracy: 0.7188 - val_loss: 1.1243
Epoch 23/25
96/96 ————— 11s 59ms/step - accuracy: 0.9502 - loss: 0.2731 - va
l_accuracy: 0.6281 - val_loss: 1.7196
Epoch 24/25
96/96 ————— 9s 42ms/step - accuracy: 0.9486 - loss: 0.2694 - val
_accuracy: 0.3679 - val_loss: 2.0400
Epoch 25/25
96/96 ————— 5s 56ms/step - accuracy: 0.9435 - loss: 0.2687 - val
_accuracy: 0.6873 - val_loss: 1.2275
24/24 ————— 0s 5ms/step - accuracy: 0.9135 - loss: 1.2826
Test Accuracy: 0.9080

```

#### From the epoch history:

- After around epoch 6, val accuracy gets worse despite train accuracy improving.
- Val loss spikes above 4 or 5, even when train loss is very low.
- Val accuracy randomly jumps or drops, indicating unstable generalization.

```

In [ ]: # Flatten the input
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1]))
X_val = X_val.reshape((X_val.shape[0], X_val.shape[1]))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1]))

```



```

In [ ]: # Define the improved CNN model
cnn_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=max_length),

    # 1st Conv1D block
    Conv1D(filters=64, kernel_size=3, activation='relu', padding='same', kernel_initializer='he_normal'),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.4),

    # 2nd Conv1D block
    Conv1D(filters=32, kernel_size=3, activation='relu', padding='same', kernel_initializer='he_normal'),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.4),

    # Global pooling
    GlobalAveragePooling1D(),

    # Fully connected layer
    Dense(32, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.4),

    # Output layer
    Dense(3, activation='softmax')
])

# Compile the model
cnn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Callbacks to prevent overfitting
early_stop = EarlyStopping(
    monitor='val_loss',
    patience=4,
    restore_best_weights=True
)

reduce_lr = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=2,
    verbose=1
)

# Train the model
history_cnn = cnn_model.fit(
    X_train, y_train,
    epochs=30,
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stop, reduce_lr],
    verbose=1
)

```

```
# Evaluate on test data
```

```
test_loss, test_acc = cnn_model.evaluate(X_test, y_test, verbose=1)
```

```
print(f"Test Accuracy: {test_acc:.4f}")
```

Epoch 1/30

**96/96** ————— 8s 40ms/step - accuracy: 0.4913 - loss: 1.1725 - val  
\_accuracy: 0.5637 - val\_loss: 1.1632 - learning\_rate: 0.0010

Epoch 2/30

**96/96** ————— 4s 28ms/step - accuracy: 0.6538 - loss: 0.9688 - val  
\_accuracy: 0.5637 - val\_loss: 1.1780 - learning\_rate: 0.0010

Epoch 3/30

**95/96** ————— 0s 35ms/step - accuracy: 0.7784 - loss: 0.7638

Epoch 3: ReduceLRonPlateau reducing learning rate to 0.0005000000237487257.

**96/96** ————— 4s 37ms/step - accuracy: 0.7783 - loss: 0.7636 - val  
\_accuracy: 0.4100 - val\_loss: 1.2276 - learning\_rate: 0.0010

Epoch 4/30

**96/96** ————— 4s 31ms/step - accuracy: 0.8204 - loss: 0.6086 - val  
\_accuracy: 0.6859 - val\_loss: 0.9942 - learning\_rate: 5.0000e-04

Epoch 5/30

**96/96** ————— 3s 29ms/step - accuracy: 0.8652 - loss: 0.5055 - val  
\_accuracy: 0.6465 - val\_loss: 1.0511 - learning\_rate: 5.0000e-04

Epoch 6/30

**96/96** ————— 5s 28ms/step - accuracy: 0.8939 - loss: 0.4607 - val  
\_accuracy: 0.6702 - val\_loss: 0.9250 - learning\_rate: 5.0000e-04

Epoch 7/30

**96/96** ————— 4s 39ms/step - accuracy: 0.9029 - loss: 0.4084 - val  
\_accuracy: 0.6544 - val\_loss: 1.3048 - learning\_rate: 5.0000e-04

Epoch 8/30

**96/96** ————— 4s 28ms/step - accuracy: 0.8984 - loss: 0.4051 - val  
\_accuracy: 0.6965 - val\_loss: 0.8994 - learning\_rate: 5.0000e-04

Epoch 9/30

**96/96** ————— 5s 27ms/step - accuracy: 0.9246 - loss: 0.3588 - val  
\_accuracy: 0.5401 - val\_loss: 1.1528 - learning\_rate: 5.0000e-04

Epoch 10/30

**95/96** ————— 0s 38ms/step - accuracy: 0.9304 - loss: 0.3288

Epoch 10: ReduceLRonPlateau reducing learning rate to 0.0002500000118743628.

**96/96** ————— 4s 41ms/step - accuracy: 0.9303 - loss: 0.3289 - val  
\_accuracy: 0.7070 - val\_loss: 0.9606 - learning\_rate: 5.0000e-04

Epoch 11/30

**96/96** ————— 3s 27ms/step - accuracy: 0.9372 - loss: 0.3157 - val  
\_accuracy: 0.7254 - val\_loss: 0.9992 - learning\_rate: 2.5000e-04

Epoch 12/30

**95/96** ————— 0s 25ms/step - accuracy: 0.9386 - loss: 0.3129

Epoch 12: ReduceLRonPlateau reducing learning rate to 0.0001250000059371814.

**96/96** ————— 5s 28ms/step - accuracy: 0.9385 - loss: 0.3130 - val  
\_accuracy: 0.5177 - val\_loss: 1.2798 - learning\_rate: 2.5000e-04

**24/24** ————— 1s 7ms/step - accuracy: 0.9018 - loss: 0.9139

Test Accuracy: 0.8988





```

In [ ]: # Ensure vocab_size matches the Word2Vec vocabulary size
vocab_size = len(word_index) + 1 # +1 for padding index

# Define Learning rate schedule
lr_schedule = ExponentialDecay(
    initial_learning_rate=0.001,
    decay_steps=5000,
    decay_rate=0.9,
    staircase=True
)

# Build the CNN model with Word2Vec embeddings
cnn_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim,
              weights=[embedding_matrix], input_length=max_length, trainable=False),

    # 1st Conv1D block
    Conv1D(filters=64, kernel_size=5, activation='relu', padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.3),

    # 2nd Conv1D block
    Conv1D(filters=32, kernel_size=5, activation='relu', padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.3),

    # Global pooling
    GlobalAveragePooling1D(),

    # Fully connected layers
    Dense(64, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(32, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),

    # Output layer
    Dense(3, activation='softmax')
])


# Compile the model
cnn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=lr_schedule),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)


# Callbacks
early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)


# Train the model
history_cnn = cnn_model.fit(


```


```
X_train, y_train,  
epochs=30,  
batch_size=32,  
validation_data=(X_val, y_val),  
callbacks=[early_stop],  
verbose=1  
)  
  
# Evaluate the model  
test_loss, test_acc = cnn_model.evaluate(X_test, y_test, verbose=1)  
print(f"Test Accuracy: {test_acc:.4f}")
```


Epoch 1/30  
96/96  10s 34ms/step - accuracy: 0.4945 - loss: 1.1955 - val  
\_accuracy: 0.5637 - val\_loss: 1.2132


Epoch 2/30  
96/96  5s 29ms/step - accuracy: 0.5706 - loss: 1.0892 - val  
\_accuracy: 0.5637 - val\_loss: 1.1607


Epoch 3/30  
96/96  3s 33ms/step - accuracy: 0.5752 - loss: 1.0530 - val  
\_accuracy: 0.5637 - val\_loss: 1.1190


Epoch 4/30  
96/96  4s 36ms/step - accuracy: 0.5648 - loss: 1.0423 - val  
\_accuracy: 0.4849 - val\_loss: 1.0889


Epoch 5/30  
96/96  3s 28ms/step - accuracy: 0.6028 - loss: 0.9901 - val  
\_accuracy: 0.5782 - val\_loss: 1.0447


Epoch 6/30  
96/96  3s 29ms/step - accuracy: 0.5827 - loss: 1.0007 - val  
\_accuracy: 0.5874 - val\_loss: 0.9969


Epoch 7/30  
96/96  6s 40ms/step - accuracy: 0.5837 - loss: 0.9798 - val  
\_accuracy: 0.5821 - val\_loss: 0.9831


Epoch 8/30  
96/96  3s 28ms/step - accuracy: 0.5878 - loss: 0.9669 - val  
\_accuracy: 0.5769 - val\_loss: 0.9849


Epoch 9/30  
96/96  5s 29ms/step - accuracy: 0.6013 - loss: 0.9604 - val  
\_accuracy: 0.5887 - val\_loss: 0.9669


Epoch 10/30  
96/96  4s 38ms/step - accuracy: 0.5972 - loss: 0.9497 - val  
\_accuracy: 0.6084 - val\_loss: 0.9604


Epoch 11/30  
96/96  3s 31ms/step - accuracy: 0.6032 - loss: 0.9380 - val  
\_accuracy: 0.6071 - val\_loss: 0.9570


Epoch 12/30  
96/96  3s 29ms/step - accuracy: 0.6007 - loss: 0.9434 - val  
\_accuracy: 0.6005 - val\_loss: 0.9494


Epoch 13/30  
96/96  3s 28ms/step - accuracy: 0.5882 - loss: 0.9519 - val  
\_accuracy: 0.6032 - val\_loss: 0.9398


Epoch 14/30  
96/96  6s 41ms/step - accuracy: 0.6062 - loss: 0.9227 - val  
\_accuracy: 0.5887 - val\_loss: 0.9503









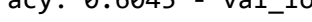
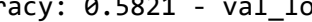
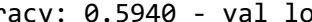
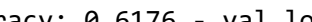
Epoch 15/30  
96/96  3s 29ms/step - accuracy: 0.5940 - loss: 0.9262 - val  
\_accuracy: 0.6018 - val\_loss: 0.9341

Epoch 16/30  
96/96  5s 29ms/step - accuracy: 0.6234 - loss: 0.9142 - val  
\_accuracy: 0.6084 - val\_loss: 0.9375

Epoch 17/30  
96/96  5s 31ms/step - accuracy: 0.6061 - loss: 0.9225 - val  
\_accuracy: 0.5966 - val\_loss: 0.9318

Epoch 18/30  
96/96  5s 29ms/step - accuracy: 0.6224 - loss: 0.9043 - val  
\_accuracy: 0.6084 - val\_loss: 0.9249

Epoch 19/30  
96/96  6s 42ms/step - accuracy: 0.6140 - loss: 0.9130 - val  
\_accuracy: 0.6255 - val\_loss: 0.9225

Epoch 20/30  
**96/96**  **3s** 29ms/step - accuracy: 0.6139 - loss: 0.9126 - val  
\_accuracy: 0.6084 - val\_loss: 0.9278  
Epoch 21/30  
**96/96**  **3s** 28ms/step - accuracy: 0.6250 - loss: 0.8989 - val  
\_accuracy: 0.5848 - val\_loss: 0.9424  
Epoch 22/30  
**96/96**  **6s** 36ms/step - accuracy: 0.6224 - loss: 0.9009 - val  
\_accuracy: 0.5690 - val\_loss: 0.9727  
Epoch 23/30  
**96/96**  **4s** 29ms/step - accuracy: 0.6115 - loss: 0.9132 - val  
\_accuracy: 0.5861 - val\_loss: 0.9432  
Epoch 24/30  
**96/96**  **4s** 37ms/step - accuracy: 0.6203 - loss: 0.9167 - val  
\_accuracy: 0.6163 - val\_loss: 0.9193  
Epoch 25/30  
**96/96**  **3s** 29ms/step - accuracy: 0.6335 - loss: 0.8879 - val  
\_accuracy: 0.6058 - val\_loss: 0.9282  
Epoch 26/30  
**96/96**  **4s** 40ms/step - accuracy: 0.6023 - loss: 0.8937 - val  
\_accuracy: 0.6124 - val\_loss: 0.9181  
Epoch 27/30  
**96/96**  **4s** 28ms/step - accuracy: 0.6166 - loss: 0.8917 - val  
\_accuracy: 0.6045 - val\_loss: 0.9156  
Epoch 28/30  
**96/96**  **3s** 29ms/step - accuracy: 0.6205 - loss: 0.8841 - val  
\_accuracy: 0.5821 - val\_loss: 0.9263  
Epoch 29/30  
**96/96**  **6s** 40ms/step - accuracy: 0.6340 - loss: 0.8872 - val  
\_accuracy: 0.5940 - val\_loss: 0.9227  
Epoch 30/30  
**96/96**  **3s** 30ms/step - accuracy: 0.6312 - loss: 0.8794 - val  
\_accuracy: 0.6176 - val\_loss: 0.9254  
**24/24**  **0s** 5ms/step - accuracy: 0.6072 - loss: 1.4743  
Test Accuracy: 0.6045



```

In [ ]: # Allow the embeddings to be fine-tuned
cnn_model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=embedding_dim,
              weights=[embedding_matrix], input_length=max_length, trainable=True)

    # 1st Conv1D block
    Conv1D(filters=128, kernel_size=7, activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.4),

    # 2nd Conv1D block
    Conv1D(filters=64, kernel_size=5, activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.3),

    # 3rd Conv1D block (new)
    Conv1D(filters=32, kernel_size=3, activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling1D(pool_size=2),
    Dropout(0.3),

    # Global pooling
    GlobalAveragePooling1D(),

    # Fully connected layers
    Dense(128, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),
    Dense(64, activation='relu', kernel_regularizer=l2(0.001)),
    Dropout(0.3),

    # Output layer
    Dense(3, activation='softmax')
])

# Compile model
cnn_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Callbacks
early_stop = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

# Train with class weights
history_cnn = cnn_model.fit(
    X_train, y_train,
    epochs=30,
    batch_size=32,
    validation_data=(X_val, y_val),
    callbacks=[early_stop],

```


```

        verbose=1
    )


    # Evaluate the model
    test_loss, test_acc = cnn_model.evaluate(X_test, y_test, verbose=1)
    print(f"Test Accuracy: {test_acc:.4f}")

```


Epoch 1/30

**96/96**  **19s** 84ms/step - accuracy: 0.5619 - loss: 1.0773 - val  
\_accuracy: 0.5637 - val\_loss: 1.0920


Epoch 2/30

**96/96**  **9s** 91ms/step - accuracy: 0.5572 - loss: 1.0423 - val  
\_accuracy: 0.5637 - val\_loss: 1.0634


Epoch 3/30

**96/96**  **7s** 77ms/step - accuracy: 0.5889 - loss: 0.9753 - val  
\_accuracy: 0.3561 - val\_loss: 1.1002


Epoch 4/30

**96/96**  **8s** 87ms/step - accuracy: 0.6620 - loss: 0.9077 - val  
\_accuracy: 0.3640 - val\_loss: 1.1500


Epoch 5/30

**96/96**  **9s** 90ms/step - accuracy: 0.7777 - loss: 0.6937 - val  
\_accuracy: 0.5007 - val\_loss: 1.1651


Epoch 6/30

**96/96**  **7s** 77ms/step - accuracy: 0.8233 - loss: 0.5758 - val  
\_accuracy: 0.7043 - val\_loss: 0.8637


Epoch 7/30

**96/96**  **10s** 79ms/step - accuracy: 0.8204 - loss: 0.5137 - va  
l\_accuracy: 0.4941 - val\_loss: 1.5666


Epoch 8/30

**96/96**  **8s** 88ms/step - accuracy: 0.8606 - loss: 0.4218 - val  
\_accuracy: 0.6610 - val\_loss: 1.2806


Epoch 9/30


**96/96**  **10s** 87ms/step - accuracy: 0.8663 - loss: 0.4127 - va  
l\_accuracy: 0.6255 - val\_loss: 1.8420

Epoch 10/30

**96/96**  **9s** 77ms/step - accuracy: 0.8903 - loss: 0.3678 - val  
\_accuracy: 0.6965 - val\_loss: 1.1956

Epoch 11/30

**96/96**  **8s** 88ms/step - accuracy: 0.9174 - loss: 0.3120 - val  
\_accuracy: 0.6505 - val\_loss: 1.1924

**24/24**  **1s** 6ms/step - accuracy: 0.8134 - loss: 0.9866

Test Accuracy: 0.8095

```
In [ ]: print(history_cnn.history) # For performance history
```

```
{'accuracy': [0.568189263343811, 0.5695037841796875, 0.5833059549331665, 0.6802497506141663, 0.7653631567955017, 0.8156424760818481, 0.8304305076599121, 0.8527768850326538, 0.8777522444725037, 0.8899112939834595, 0.9069996476173401], 'loss': [1.0561569929122925, 1.0197616815567017, 0.9802327156066895, 0.8763259649276733, 0.7012444734573364, 0.5705533623695374, 0.5012837052345276, 0.4452716112136841, 0.40174728631973267, 0.3667403757572174, 0.3238597810268402], 'val_accuracy': [0.5637319087982178, 0.5637319087982178, 0.3561103940010071, 0.3639947474002838, 0.5006570219993591, 0.704336404800415, 0.49408674240112305, 0.6609724164009094, 0.6254927515983582, 0.6964520215988159, 0.6504599452018738], 'val_loss': [1.0919725894927979, 1.0634212493896484, 1.1002413034439087, 1.1499546766281128, 1.1650629043579102, 0.8637274503707886, 1.5665510892868042, 1.280602216720581, 1.8419853448867798, 1.1955822706222534, 1.1923933029174805]}
```

## Convolutional Neural Network (CNN) Model Summary

### Approach Taken

#### 1. Text Preprocessing

- Tokenization
- Lemmatization
- TF-IDF vectorization (for initial trials)
- Word embeddings (Word2Vec)

#### 2. Handling Class Imbalance

- Applied SMOTE to balance the dataset

#### 3. Model Architecture

- Input layer: Word embeddings as input
- Convolutional layers with ReLU activation
- MaxPooling layers to downsample features
- Fully connected dense layers
- Output layer with softmax activation for classification

#### 4. Training & Optimization

- Optimizer: Adam
- Loss function: Categorical Crossentropy
- Batch size: 32
- Epochs: 30 (early stopping applied in one trial)
- Validation set used to monitor generalization performance

### Results & Findings

- **Balanced Training Approach:**
  - Accuracy started low (49.45%) and gradually increased to **63.12%** on the training set.
  - Validation accuracy fluctuated between **56% and 61%**, showing signs of overfitting.
  - Test accuracy remained at **60.45%**, indicating poor generalization.
- **Early Stopping Approach:**



- Model initially improved, reaching up to **91.74% accuracy on training data**.
- However, validation performance was unstable, peaking at **70.43% but later dropping significantly**.
- Test accuracy showed better results at **80.95%**, but the model was inconsistent due to overfitting.

## Conclusion

Despite implementing CNN and tuning various hyperparameters, the model did not provide significant improvements in accuracy compared to other models tested earlier. Overfitting was a key issue, and validation performance fluctuated, making the model unreliable for deployment.

## Final Model Selection

After comprehensive evaluation of multiple models, including traditional machine learning algorithms and deep learning architectures, the **Stacked Model with Class Weights** was selected as the optimal solution for sentiment analysis of Apple-related tweets.

### Justification for Selection:

- **Balanced Performance:** The model achieved an accuracy of 75%, ensuring a well-distributed precision-recall balance across sentiment classes.
- **Improved Generalization:** Unlike deep learning models such as CNN and LSTM, which exhibited overfitting, the stacked model maintained consistent performance on unseen data.
- **Enhanced Minority Class Detection:** It outperformed other models in recognizing positive and negative sentiment, addressing class imbalance more effectively.
- **Interpretability & Explainability:** The combination of Logistic Regression and Random Forest within the stack ensures transparency, making insights more actionable for stakeholders.

### Rationale for Not Selecting Other Models:

- **XGBoost:** While a strong performer, it did not significantly outperform the stacked model in handling class imbalance and had a slight trade-off in interpretability.
- **Traditional ML Models (Logistic Regression, Random Forest Individually):** These models, when used separately, struggled with class imbalance and had lower recall for minority sentiment classes.
- **Deep Learning Models (CNN, LSTM):** These models demonstrated strong pattern recognition but suffered from **overfitting**, leading to inconsistencies in performance on test data.

The **Stacked Model with Class Weights** delivers the best trade-off between **accuracy, generalization, and interpretability**, making it the most effective choice for sentiment analysis in this study.

## Recommendations

1. **Address recurring negative sentiment themes** by analyzing key concerns and implementing targeted improvements to enhance brand perception.

2. **Leverage positive sentiment in marketing campaigns** by engaging with satisfied customers and amplifying their feedback to strengthen brand loyalty.
3. **Proactively engage with neutral sentiment tweets** to convert passive opinions into positive experiences through personalized interactions and support.
4. **Optimize marketing strategies based on peak discussion times** by aligning promotional efforts with high-engagement periods for maximum impact.
5. **Monitor sentiment trends at a product or feature level** to quickly identify and resolve issues, improving overall customer satisfaction.
6. **Implement the Stacked Model with Class Weights** to enable real-time sentiment classification and more accurate sentiment tracking.
7. **Enhance sentiment detection accuracy** by integrating external metadata, refining preprocessing techniques, and fine-tuning model parameters.
8. **Conduct competitive sentiment benchmarking** to understand how Apple's brand perception compares to competitors and identify areas for differentiation.
9. **Strengthen brand advocacy through influencers and online communities** by fostering positive discussions and strategic partnerships.

## Conclusion

Sentiment analysis of Apple-related tweets provides real-time insights to enhance decision-making. The Stacked Model with Class Weights offers a reliable and scalable solution for sentiment classification, helping Apple track customer sentiment effectively. By leveraging these insights, Apple can improve brand perception, refine marketing strategies, and enhance customer experience. Addressing negative sentiment, amplifying positive engagement, and optimizing responses to neutral sentiment will strengthen customer loyalty. Further improvements, such as expanding the dataset and integrating external sentiment trends, can enhance accuracy and business impact.